

The Visible Hand*

Juan Sagredo[†]

October 23, 2023

Abstract

Technological innovations have allowed some sellers to collect detailed information about buyers. We study the effects of these changes in the heterogeneous precision of sellers' demand information. Precision plays a central role in the incentive to profit through more efficient trade versus more extractive trade - impacting total surplus and its distribution. The more heterogeneous buyers' preferences and the flatter sellers' costs, the greater the disposition of imprecisely informed sellers to profit through extraction despite limiting trade. Precision relaxes this tension, allowing sellers to target extraction when it is less damaging to trade. It increases aggregate surplus and benefits (hurts) buyers with a low (high) willingness to pay. However, in alternate settings, imprecisely informed sellers are predisposed to profit through trade, and expertise can reorient them toward extraction to such an extent that it decreases the surplus of all buyers along with the aggregate. In either case, precision makes sellers who acquire it more profitable, but it can benefit or hurt other sellers. The more competitive an economy and the lower the precision of sellers who upskill, the worse the profit externalities; nevertheless, externalities can be positive and allow all sellers to become more profitable. And yet, competition can also protect precision leaders by discouraging laggards from catching up.

Keywords: Asymmetric Information, Beliefs, Data, Imperfect Competition, Mechanism Design, Pricing, Search.

JEL Classification: D43, D49, D82, D83, L13.

*We are grateful for the invaluable advice of Guillermo Ordoñez, Benjamin Lester, Rakesh Vohra, and Hal Cole, as well as the helpful feedback from participants in Stanford GSB's Rising Scholars Conference, the International Industrial Organization Conference, and Penn's Theory group.

[†]Economics Department, University of Pennsylvania, 133 South 36th Street. Can be reached at sagju@sas.upenn.edu.

1 Introduction

Matt Murray (Wall Street Journal, Editor in Chief): *The perception of a lot of people is that you’ve morphed...In a lot of ways, you’re thought of as a data company more than a retail company.*

Jeff Wilke (CEO, Amazon Worldwide Consumer): *There was a corner pharmacy where I grew up. The pharmacist had been there forever. When you walked in, he knew what you liked to buy...That’s the same thing we’re doing. Our main purpose in storing your purchases is so that we can recommend something that you might want to buy the next time.*

-“Amazon’s Defense of Private Brands” (WSJ, 10/24/19).

Amazon’s well-documented harvesting and leveraging of consumer data exemplifies a sweeping transformation in how contemporary firms operate. A surge in the availability of data and in the power of analytical methods that uncover its insights¹ now helps firms better discern consumer preferences and tailor offers². These methods have spread across industries, but diffusion within each has been heterogeneous. And yet, despite the importance of this pipeline, its impact on markets is not well understood.

There are several natural questions about this technological change. First, whether it benefits us in aggregate and as consumers is of immediate importance. In this vein, policymakers have raised concerns about potentially detrimental effects on competition (FTC (2012, 2013, 2014), CEA (2015), UK Competition and Markets Authority (2021),) and proactively moved towards containing them with wide-ranging measures, such as the EU’s General Data Protection Regulation (European Commission (2012)). Second, how do firms benefit from investments in data expertise? The growing analytics gap³ between Amazons of the world and more traditional businesses suggests that investment is profitable, but in the right context. So, what is that context? And, how does this type of investment affect competitors? To address these questions, we leverage a standard search-theoretic framework of imperfect competition, featuring buyers with heterogeneous private valuations for quality, and introduce sellers with buyer valuation signals of heterogeneous precision. In particular, sellers are differentiated by an ex ante characteristic, which we refer to as a seller’s (data) *expertise*, and expertise increases signal precision. We characterize equilibria analytically, linking properties of information with properties of offers, and study the comparative statics of both expertise and market structure, documenting their effect on the nature of competition and social surplus, in aggregate terms, as well as its distribution between and within buyers and sellers.

Expertise fundamentally impacts a seller’s trade-off between profiting through more sales (trade) versus higher markups (rents). On the one hand, it allows sellers to rent-seek when this is least damaging to trade. On the other hand, it makes rent-seeking more profitable. The former channel increases trade efficiency and, by extension, aggregate surplus, while the latter decreases these. When sellers face buyers with more heterogeneous preferences or operating quality production technologies with flatter variable costs, they are prone to pursue inefficient large markup strategies under imprecise demand information, and expertise relaxes this motive, but in the alternate settings, imprecisely

¹Since 2018, the McKinsey Global Institute has conducted a yearly survey on the “State of AI”, “representing the full range of regions, industries, company sizes, functional specialties, and tenures”. Approximately half of all firms consistently report the adoption of AI in at least one business function, while the average number of functions has doubled since the first survey.

²Mikians et al. (2012, 2013), Hannak et al. (2014), Chen et al. (2015) document the pervasiveness of these methods among sellers, both large and small, while industry surveys by Deloitte (2018) and McKinsey (2023) echo these points, respectively finding widespread use of AI for personalization and that “Marketing and Sales” along with “Product/Service Development” are the most common business functions for AI applications.

³In McKinsey’s surveys, high-performing organizations are more than three times as likely to report that their data and analytics contributed at least 20% of earnings before interest and taxes.

informed sellers instead pursue efficient large sale strategies, and expertise can weaken this motive. Beyond its aggregate impact, we find that expertise is also inherently redistributive, not only shifting the share of surplus between buyers and sellers but also its distribution within each side of the market. On the demand side, high-valuation buyers - the principal victims of rent-seeking - broadly suffer from expertise, whereas low-valuation buyers - the principal victims of rationed trade - broadly benefit from it. On the supply side, expertise improves the profitability of sellers who acquire it, but its impact on other sellers who do not can be positive or negative. More expert competitors are more likely to offer more (less) competitive low (high) prices to low (high) valuation buyers; in this sense, expertise tends to decrease (increase) the profitability of other sellers' trade with low- (high-) valuation buyers. How much expertise impacts the profitability of trade with each type of buyer determines the sign of its profit externality. The more competitive the market (structure) and the lower the expertise of sellers who upskill, the more (less) that expertise promotes (relaxes) competition in low (high) matches, and the worse its profit externalities; nevertheless, these can be positive and create opportunities for all sellers to benefit, despite the privacy of signals and actions. Counterintuitively, these latter settings include cases where laggard sellers fall further behind and where they catch up to leaders.

We confirm the traditional effects of competition but also find a new - concerning - one. Competition does indeed weaken sellers' incentive to rent-seek, resulting in more efficient trade and a larger share of surplus for buyers. However, competition also lowers the returns of expertise, and when expertise is heterogeneous, this dampening affects sellers unequally. Indeed, both competition and imprecision make sellers' offers less sensitive to their information, even becoming information-insensitive (invariant in the seller's signal) at the extreme. This complementarity makes amateur profits particularly vulnerable to information-insensitivity - a state where they are unable to use and are divestitized from improving their expertise. Although we do not model expertise acquisition, this suggests the existence of a novel strategic equilibrium force behind the documented disparity in firm usage of these information technologies, and that competition exacerbates it.

Before proceeding with more detailed results, it is helpful to describe the components of our model. On one side of the market, there is a unit mass of buyers with low or high private valuations for the quality of a good. On the other side, there is a unit mass of sellers who produce this good operating a common technology with convex costs for quality. We introduce imperfect competition by randomly matching each buyer with either one or two sellers who make simultaneous take-it-or-leave-it offers, without any information about the number of competitors in the match, beyond the commonly known matching protocol. In this reduced-form formulation, we capture the full range of competitive intensity from monopoly to perfect competition by altering the probability of matching with only one seller. Our innovation is (a) introducing flexibility in the precision of sellers' demand information, and (b) then allowing this *ex ante* attribute to be heterogeneous among sellers. This is modeled by assuming that the distribution of buyer valuations is common knowledge, but that each seller in a match observes a signal about the matched buyer's valuation. These signals are essentially pairwise independent of other sellers' signals conditionally on the buyer's valuation⁴ and of precision that is increasing in the seller's expertise. For our purposes, a discretization with two relative levels of expertise suffices: less expert *amateurs* and more expert *sharks*. This allows us to study the impact of seller information in a rich fashion by varying the absolute level of amateur and shark expertise, through the precision of their signals⁵.

We introduce most of the core ideas in a simpler setting, where sellers offer an identical good of

⁴In other words, that is informative about the buyer's valuation, and only informative about the type of any competitors through its information about the buyer's valuation.

⁵We could also study this by varying the proportion of amateur and shark sellers, but the insights that we obtain by varying each experts' precision directly subsume those of the expert-proportion approach.

exogenously specified quality, so that expertise only orients pricing, and then proceed to the general case, where each seller chooses the quality of the goods it offers, so that expertise orients pricing and production. We find that information imbues equilibrium offers with several ordering properties. Recalling the timing of this incomplete information game, sellers match with buyers, observe a valuation-relevant signal, and update their valuation beliefs. At this interim stage, a seller's type is summarized by its posterior probability that the matched buyer is of low-valuation. In equilibrium, high-markup offers are more profitable in matches with high-valuation buyers, whereas high-sale offers, which share more of the gains from trade with buyers, are more profitable in matches with the low-valuation buyers. Therefore, the higher a seller's type (more weight on being in a low-valuation match), the more attractive and efficient its offers. And, expertise is intimately connected with sellers' types. For one, expertise makes it more likely that a seller observes a low (high) signal in a low- (high-) valuation match, and is of relatively high (low) type. But expertise also makes posterior valuation beliefs more extreme, because more precise signals trigger larger updates. Since sellers' weight on rent-seeking versus trade promotion is dictated by their type, sharks who observe high signals make the least attractive/highest markup/inefficient offers, while sharks who observe low ones make the most attractive/lowest markup/efficient ones - amateur offers are intermediate in these aspects. Sharks are, therefore, more profitable because (1) they are more likely to observe the signal that orients toward the right profitability lever (sales versus markups), and (2) they exercise each lever more aggressively. These equilibrium properties and the aforementioned comparative statics of market structure, along with expertise, are robust to the decision scope of sellers' demand information .ie whether it informs only pricing or also production.

However, when demand information can inform production, there are significant differences and insights. Equilibria change in form as sellers augment signal valuation information with that which can be extracted by screening menus made up of the price-quality that target each type of buyer; in other words, they perform second- and third- degree price discrimination concurrently. This difference in the format of offers also changes where we should look to track trade efficiency. When demand information only informs pricing, all trade efficiency is measured at the extensive margin (whether a buyer finds an acceptable offer), but when demand information also informs production, all trade efficiency is measured at intensive margin (how much quality a buyer obtains) as buyers always accept some seller's offer. Further, sellers' enhanced ability to extract buyers' information rents doesn't just alter the structure of equilibria, but also some of the comparative statics. On the demand side, high-valuation buyers are hurt by more efficient rent-extraction, whereas low-valuation benefit from sellers' additional flexibility in tailoring offers towards each type of buyer. On the supply side, more extractive offers are less attractive offers, so second-degree price discrimination weakens competition and makes expertise more beneficial for all sellers (individually and through its profit externalities). Lastly, we find that expertise heterogeneity becomes a central issue, as it is likely to be intrinsic to the connection between demand information and production, because by expanding the scope of information (pricing + production), we amplify the differential effect of competition on amateurs versus sharks, making the former's (latter's) offers less (more) responsive to additional expertise. As such, competition neutralizes amateur expertise - making it less hurtful for other sellers and discouraging its acquisition - and amplifies shark expertise - making it more beneficial for all sellers and encouraging its acquisition. Thus, contrary to what may be expected, competition can protect sharks' lead.

Literature Review - Our approach joins a line of work that leverages the canonical search-theoretic framework of Burdett and Judd (1983), which focused on the role of buyers' information frictions in equilibrium price dispersion, but generally provides a tractable workhorse model of imperfect competition. We begin with the analysis of economies where seller expertise does not inform production

(quality is exogenous), only pricing. These settings introduce heterogeneity among buyers, through their taste for quality, and sellers, through the precision of their (taste) information, to the canonical environment. They allow us to establish the fundamental relations between expertise, trade efficiency, and welfare purely through the pricing channel. We then remove this constraint and allow expertise to also inform production (quality is endogenous), so as to understand the impact of what is perhaps its most familiar facet for buyers - the fit of product offerings. These settings are closest to Garret⁶ et al. (2019), who also introduce buyer taste heterogeneity to the canonical environment but then focus on the role of screening. Our sellers screen, in these unconstrained settings, but they also have expertise, so they complement second- with third-degree price discrimination. This yields a more complex but still tractable model that features many of the relevant frictions and channels of these markets: imperfect competition among sellers, adverse selection of sellers by buyers whose preferences are private, and seller information of heterogeneous precision that guides production and pricing. Featuring expertise in these more general economies allows us to verify the robustness of the core relations derived exclusively via pricing, as well as obtain important insights into the influence of the product design channel on equilibrium offers, the extensive and intensive margins of trade, and the profit externalities of expertise.

Generally, we contribute to the extensive literature on price discrimination, and particularly to the branches that have been concerned with ramifications of second- and third-degree price discrimination. The latter's tradition extends back to Pigou (1920), while the former has been an active field since the seminal work of Mussa and Rosen (1978). A salient theme has been the double-edged nature of price discrimination: potentially increasing efficiency but also redistributing surplus. Our concern for the key factors that determine the direction of these effects aligns with research that has analyzed the influence of market structure⁷ and demand characteristics⁸. We contribute to this tradition by bridging second- and third-degree price discrimination and by linking new structural factors that drive aggregate outcomes. In detail, by having our sellers learn structurally, via buyers' preferences (rather than through reduced-form parameters), we are able to link its heterogeneity to the effects of price discrimination, whereas by modeling sellers' information in a rich fashion, we are able to link the level and distribution of expertise to the effects price discrimination.

Our work also forms part of a growing literature that studies both the micro-foundations of the analytics pipeline and its macroeconomic effects. At the data collection stage, Braghieri (2019), Ichihashi (2020), Hidir and Vellodi (2021), and Ali et al. (2023) focus on buyers' incentives to disclose information about their preferences, trading off the desirability of the proposed item and its price, while Bergemann et al. (2015), Elliot et al. (2021), Yang (2022), and Ichihashi (2023) focus on the ways an intermediary between buyers and firms can repack data to achieve different welfare objectives, whereas in Kehoe et al. (2020) learning goes in both directions (buyers learn about a firms' products by making purchases, and firms learn about buyer preferences through sales). These works are complementary to ours since they investigate the process of generating predictive precision, meanwhile we study a rich competitive setting that results once it has been acquired. In other words, we assume that data is both available⁹ and that some of its insights are known by some firms. We share some of the pro-welfare and pro-consumer effects of information and competition from this literature,

⁶Contemporaneously, Lester et al. (2019) study the flipside of Garret et al.'s setting - a lemons problem with an analogous matching mechanism where privately informed sellers obtain bids from uninformed buyers.

⁷See Holmes (1989), Armstrong and Vickers (2001), Rochet and Stole (1997, 2002), and Stole (2007).

⁸See Robinson (1933), Schmalensee (1981), Varian (1985), and Aguirre et al. (2010)

⁹Our buyers' randomization between sellers and non-concealment of their feature is done largely for tractability, but Turov et al. (2005), FTC (2012), and Acquisti et al. (2015) have each documented the obliviousness of shoppers to sellers' use of personal data. Nevertheless, it is important to remember that predictive skill is not just a function of data, but also of firms' ability to extract its insights - another dimension of observed heterogeneity.

but explicitly modeling the documented heterogeneity on both sides of the market allows us to parse their composition and even reverse them in some settings while considering a more general contract space with screening menus allows us to more accurately gauge the value of data for different types of sellers.

The link between data and production is analyzed in Farboodi and Veldkamp (2022) and in Eeckhout and Veldkamp (2022). Firms in both try to produce what will be most valued by consumers, but whereas the former focuses on the role that generation and the trade of data play in this process, the latter’s primary concern is that of costly risk. Our models all link data to the efficiency and diversity of offers, however, the main differences are in the pricing dimension. The previous papers feature homogeneous pricing (across sellers) under reduced-form demand frameworks, whereas we allow for and emphasize the influence of price discrimination under derived demands. This yields different efficiency, consumer welfare, and seller profit relations. From an empirical lens, Babina et al. (2022) conduct a comprehensive study of AI’s economic impact and find that this is channeled primarily through the product mix and innovation choices of firms, suggesting that our emphasis on the link between predictive skill and product offering is warranted. More specifically, their findings match our model’s predictions that expertise is generally efficiency-enhancing, despite sometimes weakening competition, and that the race for expertise is one primarily involving the most advanced firms.

The paper proceeds as follows: in Section 2, we introduce the general elements of the environment, in Section 3, we study economies where quality is exogenous, including their comparative statics, in Section 4, we endogenize quality. Particularly technical aspects are developed in the Appendix, but the main text presents them in varying degrees of detail where they reinforce the exposition.

2 General Environment

We study a two-sided market featuring heterogeneity in both the buyer and seller sides. A unit mass of buyers have single-unit demands and heterogeneous tastes for quality of the traded good.

Assumption 2.1 (Additively Separable Utility). *A buyer with a marginal value for quality θ_i obtains utility $u(q, x; \theta_i) = \theta_i q - x$ from paying a price of x to purchase a good of q quality.*

We assume prices are denominated in a numeraire good, of which buyers are endowed with large enough amount to purchase any equilibrium offer from sellers. Further, the utility of not trading is normalized to zero.

Sellers can produce a good of quality q using a common cost function for quality $\phi(q)$. Upon a successful transaction with a buyer for a contract (q, x) - a good of quality q and price x - sellers receive profits $x - \phi(q)$. The market between buyers and sellers follows an exogenously defined matching process. Following the style of Burdett and Judd (1983), a buyer matches with either one or two randomly and independently drawn sellers with the probability of a single match given by $\tilde{\rho} \in (0, 1)$. Because we assume the number of matches is only privately known by the buyer, sellers face uncertainty as to whether they are competing against a peer for a buyer’s purchase within a given match. As such, matched sellers assign not competing a probability,

$$\rho = P(\text{single match} | \text{matched}) = \frac{\tilde{\rho}}{\tilde{\rho} + 2(1 - \tilde{\rho})} \quad (2.1)$$

where ρ governs the level of match crowdedness. This parameter directly controls the intensity of competition facing sellers, which can range from the extremes of Bertrand competition ($\rho = 0$) to monopoly ($\rho = 1$).

While the ex-ante distribution of buyer valuations is common knowledge, sellers can also observe an additional signal about the particular valuation of the buyer with whom they match. Our principal innovation is to introduce heterogeneity amongst sellers in the precision of these signals. In practice, precision can improve through the acquisition of additional data or the improvement of the techniques used to draw insights from such data. We treat improvements in either as contributing to an overall level of seller *expertise*. We consider the simplest discretization of expertise, where the sellers are either less expert *amateurs* or more expert *sharks*. Their respective signal precisions are denoted

$$\alpha_e = P^e(\text{signal} = i | \text{matched buyer's valuation} = \theta_i) \quad i \in \{l, h\} \quad (2.2)$$

$$0.5 \leq \alpha_a < \alpha_s < 1 \quad (2.3)$$

In this sense, the two expertise categories of amateur and shark will remain fixed throughout the paper, even if the level of expertise that they represent, as captured by the precision of their signals α_e , differs across economies. The mass of amateurs (sharks) is $\mu(e) = \gamma$ ($\mu(s) = 1 - \gamma$). It is not essential that the level of crowdedness is interior ($\tilde{\rho} \in (0, 1)$), nor that sharks have imperfect signals (α_s), but these assumptions greatly simplify the analysis, so the corner cases will be treated separately.

Agents do not make choices before matching, so our setting is that of an incomplete information game that is solved at the interim stage, where a seller's type is its posterior probability of being in a low match, $p_l^{e,j} = P^e(\theta = \theta_l | j)$, which follows by Bayes' Rule,

$$p_l^{e,j} = \frac{P^e(j | \theta_l)p(\theta_l)}{P^e(j | \theta_l)p(\theta_l) + P^e(j | \theta_h)p(\theta_h)} = \frac{\alpha_e p(\theta_l)}{\alpha_e p(\theta_l) + (1 - \alpha_e)p(\theta_h)} \quad (2.4)$$

Having introduced the general environment, we will first analyze the role of information in a simple setting where sellers produce a good of exogenously fixed quality, so optimal offers take the form of a single quoted price per unit x . Subsequently, we will endogenize quality, allowing sellers to choose the quality of the goods they offer. This will allow us to analyze the connection between expertise and the fit of the products available to buyers but give rise to a more complex set of optimal offers in the form of screening menus composed of quality-price contracts $((q_l, x_l), (q_h, x_h))$.

3 Exogenous Quality Setting

We first consider an economy where sellers produce a homogeneous good of quality $q = 1$ at a unit cost of κ , such that there exist gains from trade with both types of buyer,

Assumption 3.1. *Gains From All Trade*

$$\kappa < \theta_l < \theta_h \quad (3.1)$$

The remaining aspects of the environment (matching process, information environment, and timing) remain exactly as specified in Section 2. In particular, a buyer matches with one or two sellers, who observe informative signals about its valuation and then extend simultaneous take-it-or-leave-it price offers. The buyer responds simply by accepting the lowest price offer, provided it is below her valuation and so entails positive utility.

By restricting to this simpler setting without product design (quality) choices, we can transparently characterize the role of seller demand information on equilibria. Sellers more convinced that they are matched with high buyers will offer higher prices, often beyond low buyers' valuation, and thus forego trade in some matches - an inefficiency. Inversely, sellers who are more convinced that they are

matched with low buyers will offer lower prices, substantially below high buyers' willingness to pay, and thus forfeit excessive rents in some matches - a profit killer. Improved information about buyers' valuations allows sellers to offer more optimal prices in each match and thus alleviate both problems. Their revised prices are acceptable, but not excessively generous. The result that information benefits individual sellers is intuitive and expected, but in its externalities, we find results that are both deep and surprising. These reveal the nuanced effects information has on buyers of each valuation, other sellers, and aggregate welfare. Before discussing these findings, we formally introduce the seller's problem and equilibria.

3.1 Seller Problem and Equilibrium Concept

The expected profits of a type $p_i^{e,j}$ seller from a price offer x are the product of its profits per sale, $x - \kappa$, times its expected sales, $P^e(x \text{ sale}|j)$. As in private value auctions, sellers know the value of having an offer accepted at the time of submission, the profits per sale $x - \kappa$, but they must infer the probability of its acceptance. We call the *sales* associated with an offer exactly that probability, which reflects the uncertainty that a seller has about four aspects: (1) the buyer's valuation, (2) whether the buyer also matched with another seller, (3) the type of the potential competing seller, and (4) the offer of such a seller. A type $p_i^{e,j}$ seller's optimal prices are therefore,

$$O^{e,j}(P^e(\cdot|j)) = \operatorname{argmax}_{x \in \mathbb{R}} P^e(x \text{ sale}|j)(x - \kappa) \quad (3.2)$$

and its strategy is a mixture over price offers, which we represent with a cumulative distribution function $F^{e,j}$. Fixing expected sales $P^e(\cdot|j)$, an optimal strategy is any distribution supported on the optimal price set $\operatorname{supp}(F^{e,j}) \subseteq O^{e,j}(P^e(\cdot|j))$.

Expected sales have quite a bit of structure in a Bayes-Nash equilibrium - our equilibrium concept. By Bayesian Rationality, the expected sales of a type $p_i^{e,j}$ seller's can be decomposed as,

$$\begin{aligned} P^e(x \text{ sale}|j) &= \rho + (1 - \rho) P^e(x \text{ sale}|\text{competitive match}, j) \\ &= \rho + (1 - \rho) \left(P^e(x \text{ sale}, \theta_l|\text{competitive match}, j) + P^e(x \text{ sale}, \theta_h|\text{competitive match}, j) \right) \\ &= \rho + (1 - \rho) \left(P^e(\theta_l|\text{competitive match}, j) P^e(x \text{ sale}|\theta_l, \text{competitive match}, j) \right. \\ &\quad \left. + P^e(\theta_h|\text{competitive match}, j) P^e(x \text{ sale}|\theta_h, \text{competitive match}, j) \right) \\ &= \rho + (1 - \rho) \left(p_l^{e,j} (1 - F(x|\theta_l)) \mathbb{1}(x \leq \theta_l) + p_h^{e,j} (1 - F(x|\theta_h)) \right) \end{aligned} \quad (3.3)$$

where $F(x|\theta_i)$ is the seller's expected distribution of competitor price offers matches with θ_i valuation buyers. This decomposition highlights and quantifies the distinct sources of uncertainty facing sellers - crowdedness (ρ terms), the buyer's valuation ($p_i^{e,j}$ terms), and competitor offers ($F(x|\theta)$ terms). The last of these encapsulates the uncertainty about the types of competitors and the price that each type would offer. The distribution of competitor offers $F(\cdot|\theta_i)$ emerges naturally in equilibrium as the

aggregation of each seller type's strategy,

$$\begin{aligned}
F(x|\theta) &= \sum_{\substack{e \in \{a,s\} \\ j \in \{l,h\}}} P(\text{expertise } e \text{ competitor}) P(\text{expertise } e \text{ seller observes } j|\theta) P(e \text{ seller offers } x|\text{signal } j) \\
&= \sum_{\substack{e \in \{a,s\} \\ j \in \{l,h\}}} \mu(e) P^e(j|\theta) F^{e,j}(x)
\end{aligned} \tag{3.4}$$

where again $\mu(e)$ is the mass of e expertise sellers, $P^e(j|\theta)$ is the probability that they observe the signal j in a match with a θ valuation buyer, and $F^{e,j}(x)$ is the cumulative distribution of prices they would offer upon observing j .

The Bayes-Nash equilibrium of the game is defined in terms of these individual and aggregate distributions.

Definition 3.1 (Bayes-Nash Equilibrium). *A Bayes-Nash equilibrium of an economy with exogenous quality is a pair of competitor price distributions, $(F(x|\theta_l), F(x|\theta_h))$, and vector of seller type strategies, $\{F^{e,j}\}_{e \in \{a,s\}, j \in \{l,h\}}$, for each type of seller $p_l^{e,j}$, such that,*

- **Rational Beliefs:** *Sellers' expected sales accurately reflect the distribution of competing prices, satisfying (3.3), and competitor price distributions accurately reflect the prices chosen by sellers, satisfying (3.4).*
- **Seller Optimality:** *Seller strategies are optimal given their expected sales, satisfying $\text{supp}(F^{e,j}) \subseteq O^{e,j}(P^e(\cdot|j))$ for all $(e,j) \in \{a,s\} \times \{l,h\}$.*

The technical properties of the equilibria lend great tractability. In particular, their component distributions are continuous, almost everywhere differentiable, and strictly increasing on at most two convex interval(s). Whereas, their qualitative properties are intuitive and revealing, such as the fact that higher prices are more profitable in high matches and are offered by sellers whose posteriors place more weight on being matched with a high valuation buyer - crisply illustrating the role of information. Proofs of these points are constructive, heavily reliant on economic reasoning, and relatively brief, so it will be insightful to go through them in the main text.

3.2 Equilibrium

3.2.1 Corner Cases: Monopoly and Perfect Competition

We first analyze the two corner cases of monopoly ($\rho = 1$) and perfect competition ($\rho = 0$) to fix ideas. These allow us to easily extract key insights that extend to interior economies with $\rho \in (0, 1)$ while also motivating their study, as we understand the limitations of conducting the analysis at the corners.

Monopoly - A monopolist's pricing choice involves the traditional trade-off between sales and profits per sale. Matched buyers have no alternative offers; they either accept the seller's offer to buy the good for a price of x and obtain utility $\theta_i - x$, or they reject it and obtain zero utility. Optimal take-it-or-leave-it offers must, therefore, give zero surplus to the lowest (valuation) buyer who would accept it, meaning that a seller's optimal offer must be equal to some buyer's valuation. A monopolist sets a sales maximizing price of $x = \theta_l$, and trades with every buyer, or a profit-per-sale maximizing price of $x = \theta_h$, and only trades with high buyers.

$$\theta_l - \kappa \leq (1 - p_l^{e,j})(\theta_h - \kappa) \iff p_l^{e,j}(\theta_l - \kappa) \leq (1 - p_l^{e,j})(\theta_h - \theta_l)$$

Which of the two prices the monopolist prefers is determined by her valuation beliefs, summarized by $p_l^{e,j}$. If she is sufficiently convinced that the buyer's valuation is low (high $p_l^{e,j}$), the low price $x = \theta_l$ is optimal, since the expected value of its additional sales, $p_l^{e,j}(\theta_l - \kappa)$, is greater than the expected cost of its discount, $(1 - p_l^{e,j})(\theta_h - \theta_l)$; otherwise, if she is sufficiently convinced that the buyer's valuation is high, she expects targeting their rents to be most profitable and chooses a price of $x = \theta_h$. The threshold probability p_l^* dictating this choice is given by,

Proposition 3.1 (Monopoly). *If $\rho = 1$, a seller of type $p_l^{e,j}$ offers a price¹⁰,*

$$x^*(p_l^{e,j}) = \begin{cases} \theta_h & \text{if } p_l^{e,j} < p_l^* \\ \theta_l & \text{otherwise} \end{cases} \quad (3.5)$$

for the threshold probability,

$$p_l^* = \frac{\theta_h - \theta_l}{\theta_h - \kappa} \quad (3.6)$$

Price dispersion is closely related to the precision of information. As the precision α_e of an e expertise seller's signal increases, its posteriors upon each signal become more dispersed ($p_l^{e,l} - p_l^{e,h}$ increases), and because monopolist prices are monotone in their type, posterior dispersion translates into price dispersion. The precision of information dictates whether the seller prefers to engage in third-degree price discrimination, offering prices contingent on its signal, the low price of $x = \theta_l$ upon observing a low signal l and vice versa for the high signal, or otherwise offer a uniform price. This switch occurs discontinuously when a seller's expertise is large enough that its posterior upon a high signal drops below the threshold probability p_l^* .

Precision can promote or impair trade and can redistribute surplus or make all agents weakly better off. Regarding sellers, the profitability of price discrimination is strictly increasing in the precision of its signals, as this makes it more likely that a monopolist offers each buyer the sale or profit per sale maximizing price. In this sense, precision always weakly benefits sellers and strictly so if it is sufficiently high to enable price discrimination. Regarding buyers, the low obtain zero surplus, but precision can impact their probability of trading, whereas high buyers always trade, but precision can impact their information rents. The sign of the effect that precision has on low buyer trade (trade efficiency) and high buyer rents (buyer surplus) depends on the degree of ex-ante adverse selection in the economy.

A completely uninformed monopolist ($\alpha_e = 0.5$) would offer a uniformly low or high price in every match. Its choice depends on the degree of adverse selection, which we quantify by the difference in profitability between offering the uniformly low and high prices. Therefore, adverse selection increases in the (a) mass of high buyers p_h and (b) the dispersion between the valuations of high and low buyers $\theta_h - \theta_l$.

Definition 3.2 (Degree of Adverse Selection). *An economy is said to have low (high) adverse selection if an uninformed monopolist would set a uniformly low (high) price.*

In economies with low adverse selection, uninformed sellers trade in every match, so information cannot improve trade efficiency. As long as signals are sufficiently imprecise, sellers also practice uniform pricing, and information does not play a role in the economy. However, when signals become sufficiently precise, sellers use them to tailor their offers. This leads them to misclassify buyers with probability $1 - \alpha_e$, offering the low price to high buyers and the high price to low buyers. The former

¹⁰We break profitability ties in favor of trade, but that is immaterial.




Table 1: Consider a setting where every seller has expertise e and buyers always match with a single seller ($\tilde{\rho} = \rho = 1$). **Low** (dashed) versus **High** valuation buyer outcomes.

pricing mistake is redistributive, converting seller rents into high buyer information rents, but the latter is inefficient, not just hurting the seller but also destroying gains from trade.

Inversely, in economies with high adverse selection, uninformed sellers only trade with high buyers, so additional information is crucial to facilitating trade with low buyers. Sellers only have a sufficient incentive to offer a low price upon a low signal, when its signals are sufficiently precise. Information and the ensuing price discrimination not only increase trade, with the additional gains from trade being captured by sellers, but it also makes allows high buyers to obtain some information rents. In this sense, additional precision can make all agents weakly better off.

Proposition 3.2 (Monopoly). *Consider an economy where sellers are monopolists in every match ($\rho = 1$).*

If adverse selection is low,

- **Sellers:** *Sellers with more precise signals are more profitable.*
- **Buyers:** *Low buyer surplus is zero, but high buyer surplus decreases in the precision of a matched seller's signal, strictly if the precision is high enough.*
- **Efficiency:** *A perfectly uninformed and informed seller trade in every match. As precision increases between these corner cases, a seller's probability of trade decreases discontinuously, at the level of precision that it starts price discriminating, and then strictly increases.*

Whereas, if adverse selection is high,

- **Sellers:** *Sellers with more precise signals are more profitable.*

- **Buyers:** Low buyer surplus is zero. High buyer surplus is zero under a perfectly uninformed or informed seller, but as the precision of a matched seller's signal increases between these corner cases, their surplus in matches with sellers of this expertise discontinuously increases, at the level of precision when they start price discriminating, and then decreases strictly.
- **Efficiency:** A seller's probability of trade in each match increases in the precision of its information, strictly if the precision is large enough.

In interior economies, the level of adverse selection and the recipient seller's precision of information are also some of the principal variables that dictate the impact of additional expertise, and their qualitative impact is analogous, but the level of crowding matters as well. As for the nature of its welfare effects in these economies, expertise's aggregate can still be a force for efficiency or inefficiency, whereas distributionally, these economies will allow us to measure the impact of expertise on other sellers - its profit externalities - and give rise to a non-trivial relation between it and the welfare of low buyers, who will be able to appropriate some of the gains from trade.

Perfect Competition - At the other end of perfect competition ($\rho = 0$), the unique equilibrium outcome is Bertrand price competition, where sellers offer the good at a price equal to their marginal cost and buyers capture all gains from trade.

Proposition 3.3 (Perfect Competition). *If $\rho = 0$, the unique Bayes-Nash equilibrium involves almost every seller offering the good at a price $x = \kappa$ with probability 1.*

Therefore, this is the wrong setting to study the role of sellers' information about buyers' willingness to pay for a vertically differentiated good, as in such a competitive environment, there is no strategic advantage from preference information. Nevertheless, the setting highlights the familiar pro-efficiency and pro-consumer surplus effects of competition, and these extend to interior economies with imperfect competition. However, in these, we will also show the heterogeneous impact that it has on sellers of different expertise.

3.2.2 Structure

By (3.4), equilibrium distributions of competitor offers in a low or high match - $F(x|\theta_l)$ and $F(x|\theta_h)$, respectively - are averages of each type of seller's strategy, weighted by the mass of sellers with the respective type in matches with buyers of the respective valuation. Since buyers of either valuation have a strictly positive probability of matching with any type of seller, these distributions share identical supports - their differences lie in how mass is allocated within these.

Proposition 3.4 (Identical Supports). *Equilibrium price distributions $(F(x|\theta_l), F(x|\theta_h))$ of competitor price offers in low and high matches have identical supports,*

$$\text{supp}(F(x|\theta_l)) = \text{supp}(F(x|\theta_h)) \quad (3.7)$$

Further, so long as some buyers obtain one offer but others obtain two, the equilibrium distributions - both aggregate price distributions and, by extension, seller strategies - are atomless.

Proposition 3.5 (Continuous Distributions). *Equilibrium distributions $F(\cdot|\theta_i)$ and $F^{e,j}(\cdot)$ are atomless.*

Sellers always have a strictly positive probability of competing against a peer of identical type, so a price greater than its cost per unit κ cannot be an atom, because any infinitesimal discount would

sacrifice negligible profits per sale in exchange for a discrete increase in sales. Additionally, any price offer equal to unit cost $x = \kappa$ is strictly dominated by some more expensive one, because the latter would generate some profits in each sale and at least generate sales in noncompetitive matches. This rules out an atom at the unit cost κ and implies that equilibrium distributions of competitor offers, $F(x|\theta)$, are continuous.

Since a loss of sales is the only deterrent to a price increase, and sales are only lost by becoming more expensive than another seller's offer or a buyer's valuation, competitor offer distributions must be locally increasing at any price in their support that is not equal to some buyer's marginal valuation.

Proposition 3.6 (Strictly Increasing Distributions). *Given prices $x, x' \in \text{supp}(F(x|\theta))$ such that $x < x' < \theta_l$ or $\theta_l < x < x' < \theta_h$, $F(x|\theta)$ is strictly increasing on the interval $[x, x']$ for $\theta \in \{\theta_l, \theta_h\}$.*

An immediate implication is that the price offers to any buyer are distributed in at most two disjoint intervals: one formed by prices weakly below the low valuation, $x \leq \theta_l$, and another formed by prices strictly above it, $\theta_l < x$. As such, we only need to find the endpoints of these intervals to fully characterize the support of prices - below the lowest price the distributions take value zero and above the highest they take value 1, naturally.

We start with the suprema of low-trade-permitting and overall prices. Since competitor offer distributions are atomless (Proposition 3.5), a seller who offers the highest overall price expects to only successfully sell in a monopoly match, so it bids accordingly, as per Equation (3.5). Comparing across sellers types, the one most willing to offer a high price at the cost of ceding all potential sales outside of monopoly matches is exactly the type of seller who places the greatest posterior weight of being matched with a high buyer. The link between expertise and posteriors,

$$p_l^{s,h} < p_l^{a,h} < p_l < p_l^{a,l} < p_l^{s,l} \quad (3.8)$$

therefore, makes sharks who observe high signals the principal candidates for offering the highest equilibrium price, and this price is equal to high buyers' valuation, whenever these types of sellers expect such an offer to be optimal in monopoly matches.

Proposition 3.7 (Highest Price). *The highest overall equilibrium price,*

$$\bar{x} = \begin{cases} \theta_h & \text{if } p_l^{s,h} \leq p_l^* \\ \theta_l & \text{otherwise} \end{cases} \quad (3.9)$$

Similarly, a seller who offers the highest low-trade-permitting price only expects to sell to a low buyer without any lower offers from competing sellers, so it chooses the highest price that every such buyer would accept, mainly $x = \theta_l$.

Proposition 3.8 (Form of Supports). *For $\theta_i \in \{\theta_l, \theta_h\}$,*

$$\text{supp}(F(x|\theta_i)) = \begin{cases} [\underline{x}, \theta_l] \cup [\hat{x}, \theta_h] & \text{if } p_h^{s,h}(\theta_h - \kappa) > \theta_l - \kappa \\ [\underline{x}, \theta_l] & \text{otherwise} \end{cases} \quad (3.10)$$

where $\theta_l < \hat{x}$ is the lowest equilibrium price offer that only allows trade with high buyers.

It is conceptually important to highlight that price offers above low buyers' valuation are separating: they are only ever accepted by high buyers and always rejected by low ones. As such, they are chosen to be optimal in high matches; in other words, sellers offering such prices anticipate that they

will only ever be accepted by high buyers and optimize conditional on that outcome. By removing considerations about being in a low or high match and focusing entirely on high trade, every type of seller faces the same incentives when choosing separating prices, so they all agree on the optimal set of separating prices. Some types just expect that such prices sacrifice too many sales in expectation, and instead choose pooling offers that also allow them to trade with low buyers.

Since every type of seller has the same criteria for optimal separating prices, the equation that pins down their equilibrium distribution is independent of sellers' type and follows from the necessary equilibrium condition of maintaining equal expected profits in high matches. We have argued that the highest price is either θ_l or θ_h , and only the latter is separating, so separating offers should be as profitable as an offer of θ_h ,

$$\rho(\theta_h - \kappa) = (\rho + (1 - \rho)(1 - F(x|\theta_h)))(x - \kappa) \quad (3.11)$$

where we implicitly apply the convexity of the separating support (Proposition 3.8).

Proposition 3.9 (Separating Prices in High Matches). *If (3.2) holds, the equilibrium distribution of competitor offers in a high match is,*

$$F(x|\theta_h) = 1 - \frac{\rho}{1 - \rho} \frac{\theta_h - x}{x - \kappa} \quad \forall \hat{x} \leq x \leq \theta_h \quad (3.12)$$

An offer of θ_h maximizes (minimizes) high (low) match profitability,

$$\begin{aligned} \Pi(x = \theta_h | \theta = \theta_l) &= 0 \\ \Pi(x = \theta_h | \theta = \theta_h) &= \theta_h - \kappa \end{aligned}$$

and separating prices preserve these points, so they can only be offered by sellers who place sufficient posterior probability on high matches. As such, there exists a threshold posterior $\hat{p}_l \leq p_l^*$ satisfying,

$$\hat{p}_h \Pi(x = \theta_h | \theta_h) = \sup_{x' \leq \theta_l} (\hat{p}_l \Pi(x = x' | \theta_l) + \hat{p}_h \Pi(x = x' | \theta_h)) \quad (3.13)$$

that dictates which types of sellers make separating offers - if any.

Proposition 3.10 (Threshold Posterior). *Given any equilibrium, there exists a threshold posterior,*

$$\hat{p}_l = \sup_{x' \leq \theta_l} \frac{\Pi(x = \theta_h | \theta_h) - \Pi(x = x' | \theta_l)}{\Pi(x = \theta_h | \theta_h) + \Pi(x = x' | \theta_l) - \Pi(x = x' | \theta_h)} \quad (3.14)$$

such that,

- **Separating Types:** Sellers of lower type, $p_l^{e,j} < \hat{p}_l$, only offer separating prices.
- **Dual Type:** Sellers of equal type, $p_l^{e,j} = \hat{p}_l$, offer separating and pooling prices.
- **Pooling Types:** Sellers of larger type, $\hat{p}_l < p_l^{e,j}$, only offer pooling prices.

This coarse grouping of prices is a manifestation of an ordering property, linking the indirect utility that sellers offer buyers with their posterior valuation assessments, which is characteristic of economies where sellers have expertise - heterogeneous or not. In particular, we will see that not only in the current setting where quality is exogenous, but also later in the general setting where quality is endogenous, the offers of sellers who place more posterior weight on low matches will be weakly more attractive to buyers of either valuation, and this relationship will often be strict. In the current

setting with exogenous quality, when any separating prices are offered, pooling prices must be strictly decreasing in the type of the seller who offers them, and this strict ordering can also be assumed to hold among separating prices without loss of generality. Ordered equilibria are tractable, intuitive, provide a crisp analytical characterization of the equilibrium structure that expertise generates, and allow us to study the comparative statics unambiguously (there is a unique ordered equilibrium). We will present the reasoning that underlies orderedness, deriving core insights about the importance of expertise and its heterogeneity in the process.

Consider the type of seller who offers the highest equilibrium pooling price $x = \theta_l$. This offer is only accepted in monopoly matches or in competitive matches where a separating price - all of which are higher - is offered, but since these prices are offered by sellers with low enough posterior p_l beliefs (below \hat{p}_l in Proposition 3.10), they are more likely in matches with *high* buyers¹¹. As such, the more a particular seller expects to face high buyers, the more it would expect to win in such matches offering the highest pooling price; in other words, a seller's expected sales from the highest pooling price decrease in its type $p_l^{e,j}$, and strictly so if any separating prices are offered in equilibrium. This is sufficient to uniquely identify the type of seller who offers the highest pooling price in the latter equilibria.

Proposition 3.11 (Lowest Pooling Seller Type). *If separating prices are offered in equilibrium, the highest pooling price $x = \theta_l$ is only offered by (a) sellers of the threshold type $p_l^{e,j} = \hat{p}_l$ or (b) sellers of the lowest type above this threshold $p_l^{e,j} > \hat{p}_l$, in which case sellers of types $\geq p_l^*$ only offer separating prices.*

However, when all equilibrium offers are pooling, there is no reason why offers should necessarily be ordered, and indeed there exists a symmetric equilibrium where all types have the same mixed strategy, so expertise does not play role. We solve for this symmetric equilibrium in the Appendix, but focus in the main text on economies where expertise is not superfluous.

Assumption 3.2 (Some Separating Offers). *Shark precision is such that some of their offers, upon observing a high signal, are separating ($p_l^{e,h} < \hat{p}_l$).*

We can obtain the distribution of competitor offers over the remaining pooling offers in economies that satisfy Assumption 3.2 through an inductive argument that extends the aforementioned expected sales logic. In particular, given an $x < \theta_l$ and the inductive assumption that there exists a p'_l such that equilibrium prices $x < x' \leq \theta_l$ are offered only by sellers of type $p_l^{e,j} \leq p'_l$, the expected sales logic implies that x' must only be offered by sellers of type p'_l or by those of the lowest successive type $p_l^{e,j} > p'_l$.

Theorem 3.1 (Pooling Price Ordering). *If (3.2) holds, for any pair of equilibrium prices $x \in \text{supp}(F^{e,j}), x' \in \text{supp}(F^{e',j'})$,*

$$x < x' \iff p_l^{e,j} \leq p_l^{e',j'} \quad (3.15)$$

It is also without loss to assume that this strict ordering extends to separating prices, since all sellers agree on the optimal separating prices and the equilibrium distribution of competitor offers in high matches is uniquely determined over these separating prices by (3.12).

Lemma 3.1 (Separating Price Ordering). *Given an equilibrium distribution of competitor prices $(F(x|\theta_l), F(x|\theta_h))$, there exists a set of optimal seller strategies $\{F^{e,j}(x)\}_{e \in \{a,s\}, j \in \{l,h\}}$*

- Equation (3.4) is satisfied.

¹¹Which are more likely produce the high signal that gives rise to low posterior p_l s among sellers

- *Separating Price Ordering: Given any two types $p_l^{e,j} < p_l^{e',j'} \leq \hat{p}_l$,*

$$[\max(\hat{x}, \underline{x}^{e,j}), \bar{x}^{e,j}] = [\hat{x}, \theta_h] \cap \text{supp}(F^{e,j}(x)) < \text{supp}(F^{e',j'}(x)) = [\bar{x}^{e,j}, \bar{x}^{e',j'}]$$

with \hat{p}_l satisfying (3.14).

Recapping, so long as any separating prices are offered in equilibrium, we can focus on equilibria where prices are strictly ordered by the seller's type: if both prices being compared are separating the ordering property is imposed without loss, whereas if at least one pools ordering is a necessary equilibrium property.

The ordering of each seller type's strategy implies that both the equilibrium distributions of competitor prices in either type of match and the profitability of prices in either type of match both have ordering properties. As for the distributions, note that the average mass of sellers from each type in a match depends on the buyer's valuation - in high (low) matches, more amateurs and sharks observe a signal of h (l) - so, when higher types offer low prices, the equilibrium distribution of competitor offers in high matches first-order stochastically dominates that in low matches (by Bayesian rationality (3.4)).

Proposition 3.12 (Competitor Bid Ordering). *If the supports of seller types' price distributions are ordered as per,*

$$\text{supp}(F^{e,l}) \leq \text{supp}(F^{e,h}) \quad \text{for } e \in \{a, s\} \quad (3.16)$$

the distribution of prices in high matches first-order stochastically dominates that in low matches,

$$F(x|\theta_h) < F(x|\theta_l) \quad \forall x \in \text{supp}(F(y|\theta)) \quad (3.17)$$

For the profitability ordering, note that separating offers are equally profitable in either type of match, but the equilibrium profitability of pooling offers is such that the types who offer them are indifferent. Necessarily then, for any two prices $x, x' \in \text{supp}(F^{e,j})$ offered by sellers of type $p_l^{e,j}$,

$$p_h^{e,j} \Pi(x|\theta_h) + p_l^{e,j} \Pi_h(x|\theta_l) = p_h^{e,j} \Pi(x'|\theta_h) + p_l^{e,j} \Pi_h(x'|\theta_l) \quad (3.18)$$

$$\frac{p_l^{e,j}}{p_h^{e,j}} = \frac{\Pi(x|\theta_h) - \Pi(x'|\theta_h)}{\Pi(x'|\theta_l) - \Pi(x|\theta_l)} \quad (3.19)$$

so cheaper pooling offers are more (less) profitable in low (high) matches. Of course, any trade-off between high and low match profits that makes a seller of type $p_l^{e,j}$ indifferent also incentivizes sellers of lower (higher) type to bid above (below), and this is why offer-to-profitability monotonicity and offer-to-seller-type monotonicity are two sides of the same coin.

Lemma 3.2 (Offer Profitability Ordering). *In an ordered equilibrium,*

- *All separating offers yield the identical average profits in low (high) matches.*
- *The average profits in low (high) matches of pooling offers increase as these become cheaper.*

This relation will be important in unpacking the profit externalities of expertise because it influences the types of offers that other sellers make and thus the kinds of ex-post regret that they face.

Lastly, we note that, when an ordered equilibrium exists, it is unique. This follows by induction, since the optimal offers of a seller are determined by the mass of bids above, which shape its expected sales, and these are ordered by the offering seller's type. The familiar base case is that of sellers with the lowest type, sharks who observe h . No type of seller bids above them, so their strategy

is uniquely determined. Inductively, given a seller type $p_l^{e',j'}$ and uniquely determined strategies of lower type sellers ($p_l^{e,j} < p_l^{e',j'}$) who offer higher prices, the strategy of a $p_l^{e',j'}$ type seller is also uniquely determined. We carry out this explicit process and derive sellers' strategies $\{F^{e,j}\}_{e \in \{a,s\}, j \in \{l,h\}}$ in the Appendix.

Theorem 3.2 (Unique Ordered Equilibrium). *There is a unique ordered equilibrium.*

3.3 Over- and Under-bidding

Sellers' offers are optimal under their interim expected sales assessments. However, since their information is imperfect, their offers in each match are not necessarily the most profitable ones ex-post. These offers are the ones that maximize profits when sales are conditioned on all information about the match - the buyer's valuation, the existence of competing sellers, and the type of any competitors.

Information about the number of competitors and their type complements that of the buyer's valuation, so we can generalize Lemma 3.2 to situations where an individual seller might have varying degrees of additional information. In particular, we consider three levels of nested information where the seller knows (1) the buyer's valuation, the presence of a competing seller, and its type, (2) the buyer's valuation and the presence of a competing seller, and (3) only the buyer's valuation.

Lemma 3.3 (Optimal Prices Conditionally on Match Type). *In the ordered equilibrium, conditionally on the buyer's valuation, the number of competitors, the type of the competitor,*

- *If the buyer's valuation is high, the optimal competitive price is*
 $X^*(\theta_h, \text{competitive}, p_l^{e,j} \text{ type competitor}) = \underline{x}^{e,j}.$
- *If the buyer's valuation is low, the optimal competitive price is*
 $X^*(\theta_l, \text{competitive}, p_l^{e,j} \text{ type competitor}) = \min(\underline{x}^{e,j}, \theta_l).$

whereas, conditionally on only the buyer's valuation and the number of competitors,

- *If the buyer's valuation is high, the optimal price is* $X^*(\theta_h, \text{competitive}) = \hat{x}.$
- *If the buyer's valuation is low, the optimal price is* $X^*(\theta_l, \text{competitive}) = \underline{x}.$

and, conditionally on only the buyer's valuation,

- *If the buyer's valuation is high, the set of optimal price offers is* $X^*(\theta_h) = [\hat{x}, \theta_h].$
- *If the buyer's valuation is low, the optimal price is* $X^*(\theta_l) = \underline{x}.$

Whereas, if the seller knew that it has no competitors and the buyer's valuation, then its optimal price offer is the valuation.

Since equilibrium distributions of prices are atomless (by Proposition 3.5), almost every offer that sellers make is suboptimal under the perfect information posterior. Naturally, if they sell, their offer in that particular match was too low (over-bid buyer surplus), and inversely if they do not (under-bid buyer surplus), but much of this regret is unavoidable. Indeed, no seller has information about the number of competitors, and even knowing a competitor's type, there would remain uncertainty about its offer, since equilibrium strategies are mixed. However, a part of regret is linked to the buyer's valuation, as winning (losing) makes having been in a high (low) match more likely, and sellers with more expertise have more precise valuation information. The right measure of ex-post regret, therefore, compares a seller's offer with the optimal one it would have made if it had known the

buyer's valuation. Under this metric, sharks' more precise information brings their posterior closer on average to the one conditioned on the buyer's valuation and allows them to obtain higher profits/lower regret.

To connect with our analysis of economies with endogenous quality choices, it is helpful to take a closer look at the connection between an offer's type (separating vs. pooling) and its vulnerability to each type of regret (under- versus over-bidding regret). Under-bidding affects both separating and pooling offers, as a loss indicates having been in a low match and offered too high a price ($x > X^*(\theta_l)$). However, over-bidding only affects pooling offers, as a win indicates having been in a high match and separating offers are optimal in these, but pooling offers are not ($x < X^*(\theta_h)$). Separating offers, therefore, allow sellers to minimize their over-bidding and will be the preferred mechanism in economies with endogenous quality choices, where sellers can screen even more profitably by offering menus of price-quality contracts $((q_l, x_l), (q_h, x_h))$.

3.4 Comparative Statics

With a clear understanding of the structure and analytical form of ordered equilibria, we study the effect of crowdedness and expertise on prices, trade efficiency, and the level as well as the distribution of aggregate surplus between and within buyers and sellers.

In the case of buyer welfare, we examine the average utility that buyers of each valuation obtain,

$$\mathcal{W}_b(\theta) = \rho E[u(1, \min(\theta, x_i); \theta)] + (1 - \rho) E[\max_{i \in \{1, 2\}} u(1, \min(\theta, x_i); \theta)] \quad (3.20)$$

where x_i are iid draws from the conditional equilibrium marginals $F_i(x_i|\theta_i)$. Seller profits from low and high buyer matches depend on (1) the average number of matches they are in, (2) their sales per match, and (3) their profits per sale. Since sellers with equal expertise are otherwise identical, we track average profits in low and high matches conditioning on expertise,

$$\mathcal{M} = \tilde{\rho} + 2(1 - \tilde{\rho}) \quad (3.21)$$

$$\Pi^e = \left[\underbrace{\mathcal{M}}_{\text{number of matches}} \underbrace{P^e(x \text{ sale}, j)}_{\text{probability of selling}} \underbrace{x - \kappa}_{\text{profits-per-sale}} \right] e \quad (3.22)$$

where we average over the probabilities that the seller observes each signal $j \in \{l, h\}$ in a match with a buyer of each valuation $\theta_i \in \{\theta_l, \theta_h\}$, and the prices $x \sim F^{e,j}$ when they are of either interim type $p_l^{e,j}$. Finally, because inefficient trade occurs solely from low buyers being priced out, we track efficiency by the probability low buyers trade in any given match,

$$\mathcal{Q} = \rho F(x \leq \theta_l | \theta_l) + (1 - \rho) P_{x_1, x_2 \sim F(\cdot | \theta_l)}(\min(x_1, x_2) \leq \theta_l | \theta_l) \quad (3.23)$$

Crowding promotes competition, leading to the typical effects of greater efficiency and buyer surplus. Sellers, anticipating “more” competitors in the average match, offer more attractive terms to buyers to sustain optimal sales. As sellers decrease prices across the board, low buyers have a higher probability of receiving an acceptable offer, thus boosting trade efficiency and the aggregate surplus, all of which is captured by buyers.

The effects of variation in precision are more nuanced, depending on both the level of adverse selection and crowdedness. For buyers, the low typically benefit from increased precision, while the high are harmed by it. Generally speaking, this difference stems from the upside unique to low buyers, mainly that precision reduces incidents of misclassification where sellers observe a high signal

and follow through with more expensive offers. However, in a range of settings, these general trends can be reversed. The combination of low adverse selection and low crowding can make precision harmful for all buyers, because sales are relatively inelastic in these economies, so precision weakly incentivizes sales upon a low signal but strongly incentivizes extraction upon a high one. Alternatively, the combination of high adverse selection and low crowding make precision beneficial for all buyers, because then low signals give sellers the strong incentive needed to break away from extraction (given the inelasticity of sales and high level of adverse selection). Trade efficiency tracks low buyers' utility, so these points extend to its relation with precision with the caveat that the link between low levels of adverse selection and the inefficiency of precision is even more resilient - present even under high levels of crowding.

For sellers, precision always benefits individual sellers, while typically harming their competitors. As in the case of buyer surplus and efficiency, however, there are settings where the profit externality becomes positive, so that all sellers benefit from precision. Economies with low adverse selection and low crowding or high adverse selection and high crowding are prime candidates for precision to benefit all sellers, since both lead sellers with imprecise signals to primarily offer low prices. Precision greatly improves their extraction, relaxing the competition that peers expect from these sellers who observe a high signal.

3.4.1 Crowdedness

When buyers obtain more offers, lower ρ , sellers expect them to have more alternatives, particularly more alternatives below any price. Given a fixed set of strategies, this reduces the expected sales of sellers at any price. Their best response is to restore sales, achieved by this is done by shifting mass towards lower prices in their mixed strategies. As such, the general equilibrium effect on seller's strategies reinforces the partial equilibrium, both resulting in lower prices for buyers.

We take a closer look at this dynamic in the more relevant set of equilibria - those in which expertise separating prices are offered, so expertise matters. In the ordered equilibrium, the highest overall price \bar{x} does not depend on the level of crowding - only on the posterior of sharks that observe a high signal (by (3.7)). However, the probability of prices below depends closely on crowding. Inspecting (A.4), we note that the strategy of these sharks increases strictly (in the first-order stochastic dominance ordering) with crowding,

$$F^{s,h}(x; \rho_2) < F^{s,h}(x; \rho_1) \quad \rho_1 < \rho_2, \quad \forall x < \bar{x}$$

In other words, when there is more crowding (ρ increases), they place more mass on lower prices. Induction then allows us to extend this ordering claim to other types of sellers. The support of their prices does not just extend downwards, however, but instead shifts as their most expensive offer also becomes cheaper¹²

As we have argued, offers help buyers in two ways. Fixing seller strategies, the average minimum price offer decreases in the average number of offers. Indirectly, seller strategies decrease and shift mass towards lower prices, compounding the drop in the average minimum price offer. This benefits all buyers. High buyers obtain more surplus per trade, whereas low buyers obtain more surplus

¹²In the base case of sharks who observe h , we establish the first order stochastic dominance ordering and that $\underline{x}^{s,h}$ is strictly decreasing in $1 - \rho$. This allows us to assume by induction that $\underline{x}_l^{e,j}$ is strictly decreasing in $1 - \rho$. Since the highest price offered by a type $p_l^{e,j}$ seller, $\bar{x}_l^{e,j}$, is weakly decreasing in the lowest price offered by the adjacent type below, the inductive step implies $\bar{x}_l^{e,j}$ is weakly decreasing in $1 - \rho$. We pair this point with the fact that $P^e(\bar{x}_l^{e,j} \text{ sale} | \theta_h)$ and $P^e(\bar{x}_l^{e,j} \text{ sale} | j)$ are strictly decreasing in $1 - \rho$ (by (A.6) and (A.11)) to establish that $F^{e,j}(x; \rho_2)$ strictly first-order stochastically dominates $F^{e,j}(x; \rho_1)$, for $\rho_1 < \rho_2$, and that $\underline{x}_l^{e,j}$ is strictly decreasing in $1 - \rho$.

per trade and possibly more trades. This latter efficiency gain occurs whenever some pooling and separating prices are offered, as crowding shifts some of the mass from separating prices to pooling ones. Therefore, efficiency is weakly (but often strictly) increasing in crowdedness.

However, sellers' average profits drop, in matches with either type of buyer, and the reasoning is straightforward. Recall that the sales a matched seller expects from an x offer are,

$$P^e(x \text{ sale} | j) = \rho(1 - \mathbb{1}(x \leq \theta_l)) + (1 - \rho) \sum_{\substack{j' \in \{l, h\}, \\ e' \in \{a, s\}}} \mu(e') \left(p_l^{e, j} P^{e'}(j' | \theta_l) \mathbb{1}(x \leq \theta_l) + p_h^{e, j} P^{e'}(j' | \theta_h) \right) (1 - F^{e', j'}(x)) \mathbb{1}(p_l^{e', j'} \leq p_l^{e, j})$$

and these decrease because competitive matches because more likely (ρ and $1 - \rho$ terms) and because the distribution of competitor bids shifts mass towards lower prices ($F(x | \theta)$ terms). Therefore, any offer is less profitable in any match,

$$\Pi^{e, j}(x'; \rho) = P^e(x' \text{ sale} | j; \rho)(x' - \kappa) \quad (3.24)$$

$$\Pi^{e, j}(x; \rho_1) < \Pi^{e, j}(x; \rho_2) \leq \Pi^{e, j}(x'; \rho_2) \quad x \in \text{supp}(F^{e, j}(x; \rho_1)) \text{ , } x' \in \text{supp}(F^{e, j}(x; \rho_2)) \quad (3.25)$$

On the one hand, when matches are more crowded, they are less profitable for sellers. On the other hand, crowdedness allows them to match with more buyers, and this is particularly beneficial for sellers whose offers have a higher probability of being chosen, mainly types who offer lower prices. Nevertheless, ex-ante profits involve computing the product of the average number of matches, $M(\tilde{\rho}) = \tilde{\rho} + 2(1 - \tilde{\rho})$, with a matched seller's expected sales per match, and the conditioning on matching in these places the number of matches in the denominator throughout, via $\rho = \frac{\tilde{\rho}}{M(\tilde{\rho})}$ and $1 - \rho = \frac{2(1 - \tilde{\rho})}{M(\tilde{\rho})}$, so the average number of matches term cancel,

$$\begin{aligned} & M(\rho) \Pi^{e, j}(\rho) \\ &= E_{x \sim F^{e, j}} \left[\tilde{\rho}(1 - \mathbb{1}(x \leq \theta_l)) \right. \\ & \quad \left. + 2 * (1 - \tilde{\rho}) \sum_{\substack{j' \in \{l, h\}, \\ e' \in \{a, s\}}} \mu(e') \left(p_l^{e, j} P^{e'}(j' | \theta_l) \mathbb{1}(x \leq \theta_l) + p_h^{e, j} P^{e'}(j' | \theta_h) \right) (1 - F^{e', j'}(x)) \mathbb{1}(p_l^{e', j'} \leq p_l^{e, j}) \right] \end{aligned}$$

and average ex-ante profits satisfy an analogous set of inequalities as profits per match at any price (3.25),

$$M(\rho_1) \Pi^{e, j}(\rho_1) < M(\rho_2) \Pi^{e, j}(\rho_2) \quad (3.26)$$

by the same reasoning as before¹³. The additional sales that some types achieve are not insufficient compensation for the lower prices it takes to get them, so the profits of every type of seller decrease.

Theorem 3.3 (ρ Comparative Statics). *Consider two economies that differ only in the level of crowding $\rho_1 < \rho_2$ with the corresponding ordered equilibria,*

$$\{F(x | \theta; \rho_i), F^{e, j}(x; \rho_i)\}_{\substack{\theta \in \{\theta_l, \theta_h\} \\ e \in \{a, s\} \\ j \in \{l, h\}}}$$

¹³ $\tilde{\rho}$ decreases if and only if ρ does, so $F^{e, j}(x; \rho)$ increases if and only if $\tilde{\rho}$ decreases. As such, the terms inside the expectation, once we cancel the average number of matches, are pointwise (in x) increasing in ρ .

As crowdedness increases,

- **Distributions Strictly Increasing:** $F(x|\theta; \rho_2) < F(x|\theta; \rho_1)$ and $F^{e,j}(x; \rho_2) < F^{e,j}(x; \rho_1)$ for all $(e, j, x) \neq (s, h, \bar{x})$.
- **Supports Increasing:** $\underline{x}^{e,j}(\rho_1) < \underline{x}^{e,j}(\rho_2)$ and $\bar{x}^{e,j}(\rho_1) \leq \bar{x}^{e,j}(\rho_2)$ for all e, j .
- **Buyer Surplus Strictly Increasing:** $\mathcal{W}_b(\theta; \rho_2) < \mathcal{W}_b(\theta; \rho_1) \forall \theta$.
- **Trade Efficiency Increasing:** $\mathcal{Q}(\rho_2) \leq \mathcal{Q}(\rho_1)$.
- **Profits Strictly Decreasing:** $\Pi^e(\rho_1) < \Pi^e(\rho_2) \quad \forall e$.

We illustrate these effects in ?? and ??.




Figure 1: Buyer Surplus and Efficiency Progression The common parameters are $[\theta_l, \theta_h, p_l] = [1, 3, 0.5]$ for buyers, $\kappa = 0.5$ for sellers' unit cost, $[\alpha_a, \alpha_s] = [0.55, 0.95]$ for sellers' precision, and $\gamma = 0.5$ for the proportion of amateurs.




Figure 2: Profitability Progression The common parameters are $[\theta_l, \theta_h, p(\theta_l)] = [1, 3, 0.5]$ for buyers, $\kappa = 0.5$ for sellers' unit cost, $[\alpha_a, \alpha_s] = [0.55, 0.95]$ for sellers' precision, and $p(a) = 0.5$ for the proportion of amateurs.

3.4.2 Expertise

Crowdedness is the traditional channel that heightens competition; however, it is not the only one. The precision of a seller’s demand information also affects the distribution of prices that they extend to buyers and so the degree of competition that peer sellers expect.

Precision of information about buyers’ valuations shapes a seller’s distribution of posteriors, hence types, in two ways. First, it changes how often the seller’s signal correctly classifies the buyer; in other words, how often the seller’s type is relatively low or high in matches with a buyer of each valuation - the *classification effect*. Second, it changes how extreme its posterior valuation beliefs are upon each signal; in other words, how large its type is upon each signal - the high and low signal *type effect*. Jointly, these increase the precision of its expected sales, which are central to the profitability of its offers. The closer a seller’s expected sales are to the perfect valuation information forecast $P(x \text{ sale}|\theta)$, the more precise her sales vs. markup trade-off. In particular, a seller with greater confidence about being in a high match expects more sales at any price (in the ordered equilibrium), so it offers higher prices - and vice versa for a seller with greater confidence about being in a low match. Said differently, precision makes the offers of sellers who observe high (low) signals less (more) competitive. The strategic response of other sellers to this, the *competition effect*, is why changes in the precision of a given class of experts’ signals can also produce changes in the offers of sellers with differing expertise. These three effects of precision do not always share the same sign and prevent unambiguous statements about the impact of precision on trade efficiency and welfare. However, the model allows us to categorize its aggregate impact along intuitive structural dimensions - the level of crowdedness, adverse selection, and precision - and to concisely explain the mechanisms behind it.

The advancement of shark and amateur expertise is quite different. Outside our model, we like to think of the former as firms like Amazon and OpenAI pushing the frontier of artificial intelligence and big data methods, whereas the latter is the diffusion of data analytics practices to the wider industry landscape. We will study these two types of technological progress by varying the precision of the signals that *all*¹⁴ sharks or amateurs observe.

In general, precision tends to improve trade efficiency, benefit (hurt) low (high) buyers, and increase (decrease) the profits of sellers who (do not) obtain it. We find these results fall in line with intuition and also attempt to explain some exceptions that can arise.

Buyer Surplus and Efficiency The existence of pooling prices does not just benefit high buyers, who can then obtain information rents, but also low buyers, as competition among pooling offers intensifies in the mass and the type of sellers offering them. Therefore, buyers benefit from precision to the degree that it drives sellers to offer lower pooling prices, and economies with high adverse selection are prime candidates for this effect. In these economies, crowdedness would be the only incentive for an uninformed seller to trade with low buyers, however, as the precision of its information increases, low signals become a sufficient incentive for low trade. Low valuation buyers are able to obtain more attractive offers from these sellers, because of the classification effect (more observe low signals in low matches), low signal type effect (higher type upon a low signal), and low signal competition effects (stronger competition that these experts who observe low signals exert on any higher expertise sellers who also observe low signals). High valuation buyers can benefit as well when the low signal type and classification effects are strong enough, as sellers still misclassify them at times, but when precision is high enough, the classification effect, which hurts them, is dominant. The principal difference in

¹⁴The “effect” of changing an individual seller’s precision of information will be obvious once we understand the effects of changing that of an entire class of sellers, where effect is in quotes because although discussion about the impact on an individual seller’s actions is well-defined, there is no aggregate equilibrium impact to perturbations of measure zero agents.

economies with low adverse selection is that an uninformed seller would always offer pooling prices, preferring to trade with all buyers, so precision, which encourages higher separating prices upon a high signal, hurts high valuation buyers and can only benefit the low under enough precision so that the classification and competition effects are strong enough.

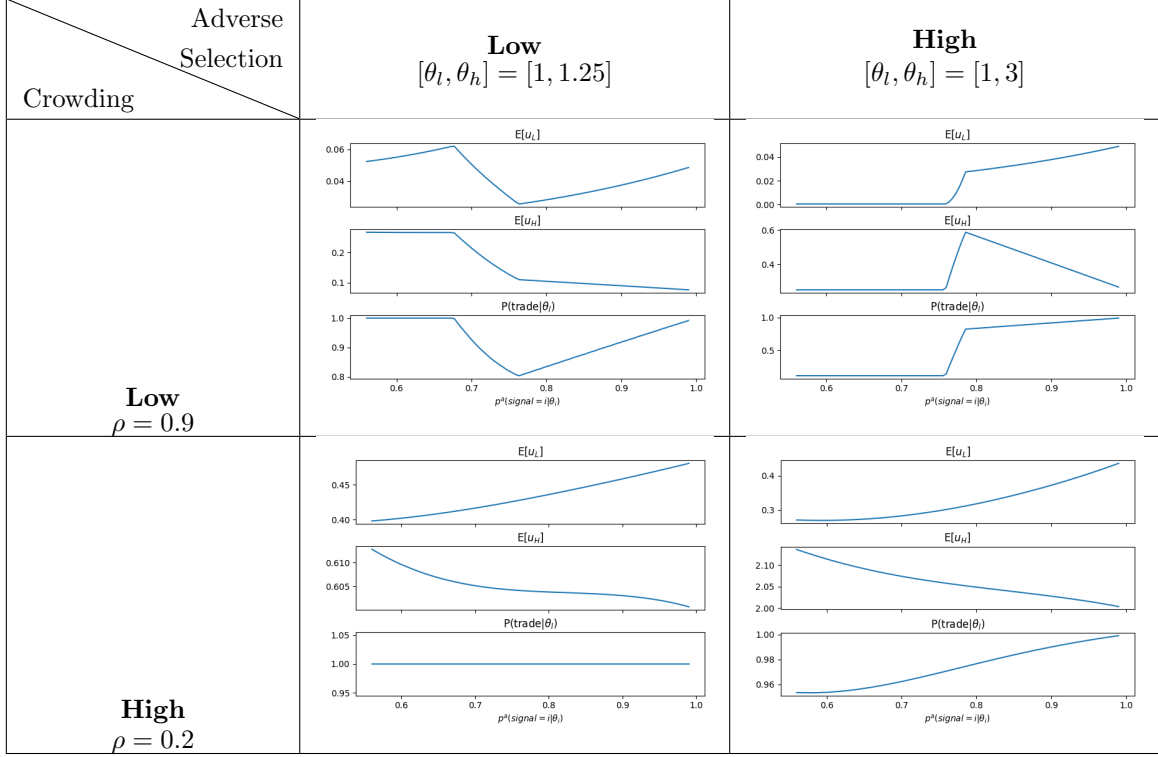


Table 2: Buyer Surplus and Efficiency Effects of Precision Almost all sellers are amateurs (90%) and we increase the precision of their signals along the respective x-axes. Each figure is representative of the effect that precision has on buyer surplus and trade at different levels of crowding and adverse selection. The common parameters are $p(\theta_l) = 0.5$ for the mass of low buyers, $\kappa = 0.5$ for sellers' unit cost, and $\alpha_s = 0.999$ for shark precision.

Efficiency increases with precision under high adverse selection, whereas under low adverse selection it is initially decreasing, but then recovers. A simple application of the product rule shows us why. Consider an economy with only one class of experts, where sellers start to offer separating prices upon a high signal but continue to pool upon a low one,

$$\begin{aligned}
 P(\text{trade}|\theta_l; \alpha) &= \underbrace{\alpha P(\text{pooling offer}|l; \alpha)}_{=1} + (1 - \alpha) P(\text{pooling offer}|h; \alpha) \\
 \frac{\partial}{\partial \alpha} P(\text{trade}|\theta_l; \alpha) &= 1 - P(\text{pooling offer}|h; \alpha) + (1 - \alpha) \underbrace{\frac{\partial}{\partial \alpha} P(\text{pooling offer}|h; \alpha)}_{<0}
 \end{aligned}$$

when precision is still too low, the increase in separating offers upon high signals is dominant, whereas if precision is higher, the decrease in misclassification is.

Profits Sellers benefit individually from more precise information, as it allows them to make more profitable offers on average - closer to the optimal prices conditionally on the matched buyer's valuation. However, since precision impacts the offers of a given seller, it also impacts the profitability of other sellers. When we consider a change in the precision of an entire class of experts (amateurs or sharks), the net response of the average amateur or shark profit is the average of their individual

benefit + profit externalities. The individual benefit is straightforward, but the profit externality brings together several components that demand closer attention.

The profit externalities of precision are clearest among sellers whose precision remains constant. First, consider the peer effect of amateur precision on sharks. If sharks observe a high signal, they always offer higher prices than amateurs, whereas if they observe a low one, they always offer lower prices than amateurs; as such, sharks' expected sales are invariant in amateur precision. In this sense, shark over-/under- bidding and the profitability of their high signal offers are all invariant in amateur precision. Amateur precision uniquely affects sharks only through the markups that they obtain upon a low signal, which occur primarily in low matches. Amateur precision increases the mass of amateurs that observe low signals in low matches and makes them keener on offering low prices upon it, so more precisely informed amateurs tend to compete more aggressively in matches with low valuation buyers, lowering sharks' profits. However, in economies with low adverse selection and low amateur precision, the stronger incentive to offer high prices upon a high signal can dominate and make amateurs compete less aggressively, even in matches with low valuation buyers, increasing sharks' profits.

The fundamental difference with shark precision is that it is the one of best-informed sellers - expertise leaders - and it is their priorities which "select" the type of equilibrium. This is why, when their information is imprecise, low adverse selection equilibria feature only pooling offers that maximize trade, and sellers' strategies are symmetric. Once sharks' information becomes sufficiently precise, they start making separating offers upon high signal; in other words, they start using their expertise. This turns on the profit externalities of expertise, which are at first discretely negative, as a switch to an ordered equilibrium is a switch to an equilibrium where sellers over-bid (buyer surplus) in high matches and under-bid in low matches - both of which decrease profitability. Nevertheless, in separating equilibria, the qualitative effects of shark precision on amateurs are qualitatively analogous to amateurs' on sharks - beneficial/costly on net under the similar combinations of crowding, adverse selection, and precision - even if the channels that bring it about differ. Amateurs offer lower prices than sharks who observe high signals and high prices than sharks who observe low ones, and we recall that the profitability of separating offers is invariant in sellers' type (all agree on the optimal separating offers), so the principal benefit that amateurs obtain from shark precision is an increase in the mass who offer separating prices during high valuation matches, whereas the principal cost is an increase in the mass who undercut them in low valuation matches.

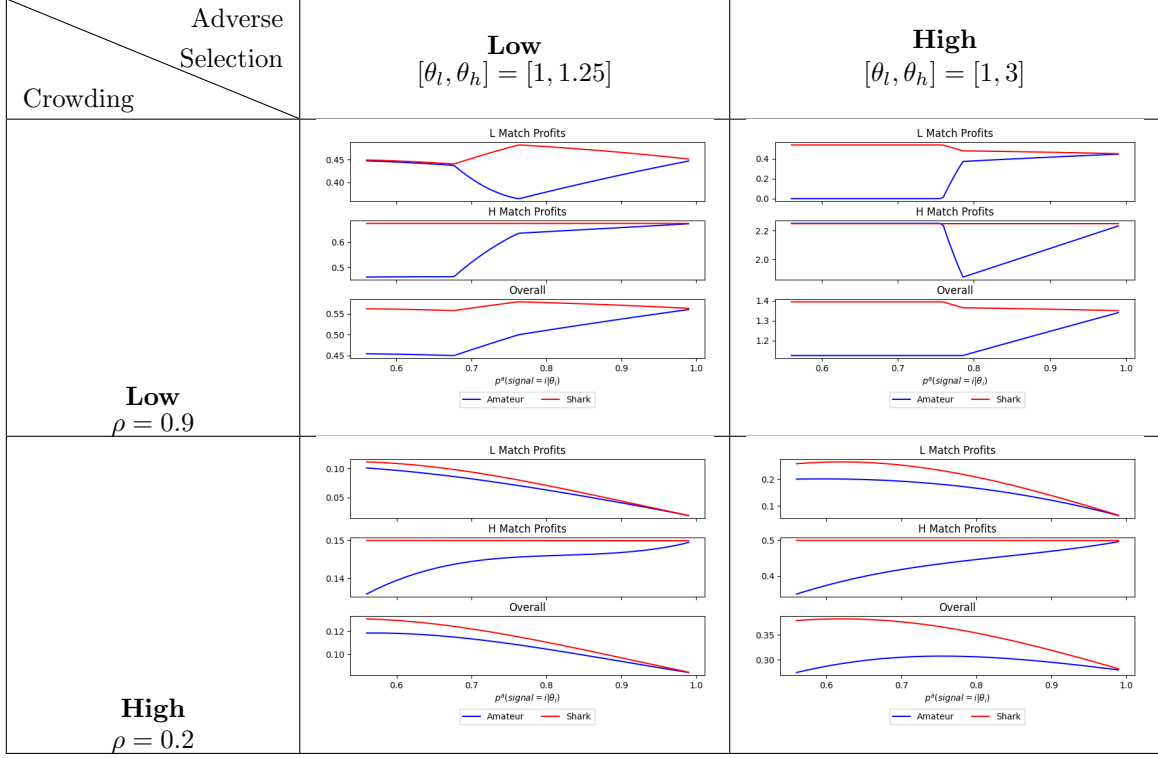


Table 3: Profits Effects of Amateur Precision A supermajority of sellers are amateurs (90%) and we increase the precision of their signals along the respective x-axes. Each figure is representative of the effect that precision has on their and sharks' profits in low matches, high matches, and the average match. The common parameters are $p(\theta_l) = 0.5$ for the mass of low buyers, $\kappa = 0.5$ for sellers' unit cost, and $\alpha_s = 0.999$ for shark precision.




Table 4: Profit Effects of Shark Precision A supermajority of sellers are sharks (99%) and we increase the precision of their signals along the respective x-axes. Each figure is representative of the effect that precision has on their and amateurs’ profits in low matches, high matches, and the average match. The common parameters are $p(\theta_l) = 0.5$ for the mass of low buyers, $\kappa = 0.5$ for sellers’ unit cost, and $\alpha_a = 0.55$ for shark precision.

4 Endogenous Quality Setting

Although pricing is one of the main use cases of demand information, another that is at least as important and particularly prevalent in sellers’ recent data analytics applications is product choice - the problem of choosing what to offer to each buyer. Abstracting away from this problem has been convenient analytically and even practically negligible in situations where buyers perceive goods as highly substitutable or sellers’ production is inflexible. However, to study settings where product differentiation and design are significant aspects, we will incorporate product choice by allowing sellers to pick the quality of the goods that they offer buyers. Although many of the core insights from exogenous product choice economies will generalize, some in exact form and others with close analogs, the significant differences lend nuance to the effects on expertise welfare and efficiency.

4.1 Seller Problem and Equilibrium Concept

The interesting cases of endogenous quality are those where the efficient quality of trade with low and high valuation buyers is different. This requires some cost convexity. We will work with a piecewise-linear cost function that kinks at the efficient qualities of trade with low- and high-valuation buyers.

Assumption 4.1 (Piecewise Linear Costs). *The cost function $\phi(\cdot)$ is piecewise linear, convex, strictly*

increasing with

$$\phi(q) = \begin{cases} \kappa_l q & q \leq q_l^* \\ \kappa_l q_l^* + \kappa_m q & q_l^* < q \leq q_h^* \\ \kappa_l q_l^* + \kappa_m (q_h^* - q_l^*) + \kappa_h q & q_h^* < q \end{cases}$$

where $0 < \kappa_l < \theta_l$, $\theta_l < \kappa_m < \theta_h$, and $\theta_h < \kappa_h$ and the efficient qualities of trade with a buyer of each valuation are given by,

$$q_i^* = \operatorname{argmax}_q \theta_i q - \phi(q) \quad (4.1)$$

Piecewise linearity makes the marginal cost of quality revisions locally constant, lending considerable tractability¹⁵. Since sellers can choose quality and its valuation among buyers is heterogeneous, they have the ability to improve their screening of buyers by going beyond single quality-price offers and instead offering each buyer a menu $((q_l, x_l), (q_h, x_h))$ composed of a pair of quality-price contracts, where the contract (q_i, x_i) is intended for a buyer of valuation θ_i . When both contracts are identical, the menu is equivalent to a single contract offer, but we will see that such a menu is never optimal for sellers, the first sign that the exogenous product choice assumption imposes serious economic limitations.

There are three components to the profits that a matched seller expects from a menu: (1) the profits per sale from each contract, (2) the probability that a buyer of a given valuation chooses each contract, and (3) the probabilities that the buyer has each valuation. Profits per sale from a contract, $\pi(q_i, x_i) = x_i - \phi(q_i)$, are simply the difference between the lump sum price x_i and the seller's cost of producing the quality q_i . The probability that a buyer of a given valuation chooses a contract is the probability that it has no better offers. By the Revelation Principle, it is sufficient to restrict attention to individually rational and incentive-compatible menus, so the probability that a buyer of valuation θ_i chooses the contract (q_j, x_j) is zero if $i \neq j$, and otherwise equal to the sum of the probability that the seller's offer is the only one¹⁶ (ρ) plus the probability that it has an inferior offer from another seller. To compute the probability that the seller's offer beats that of a competitor, we define the marginal distribution $F_i(u_i|\theta_i)$ over indirect utilities offered to θ_i valuation buyers by sellers of each type via the joint distribution of utilities in such matches,

$$F(u_l, u_h|\theta_i) = \sum_{\substack{e \in \{a, s\} \\ j \in \{l, h\}}} \mu(e) P^e(j|\theta_i) F^{e,j}(u_l, u_h) \quad (4.2)$$

where $\mu(e)$ is the mass of sellers with expertise e , $P^e(j|\theta_i)$ is the proportion of them that would observe j signals when matched with a θ_i valuation buyer (and so that would be of type $p_l^{e,j}$), and $F^{e,j}(u_l, u_h)$ is the joint distribution over indirect utility offers implied the menus that sellers of type $p_l^{e,j}$ mix over. The marginal distributions $F_i(u_i|\theta_i)$ follow directly from this joint distribution. Therefore, the probability that a θ_i valuation buyer chooses a contract (q_i, x_i) contract is $\Psi_i(u_i; F) = \rho + (1 - \rho)F_i(u_i|\theta_i)$, where $u_i = u(q_i, x_i; \theta_i)$ is the indirect utility it offers θ_i valuation buyers. Lastly, a matched seller's probability that the buyer has low and high valuation is given by its type $p_l^{e,j}$ and the respective complementary probability $1 - p_l^{e,j}$. Since each contract (q_i, x_i) determines the seller's expected profits in a match with each kind of buyer, the expected profits from the menu are therefore

¹⁵Our main results are not dependent on this restriction.

¹⁶The probability that the seller is a monopolist given that it is matched.

given by the average of these,

$$\begin{aligned}\Pi^{e,j}(q_l, x_l, q_h, x_h) &= p_l^{e,j} \Psi_l(u(q_l, x_l; \theta_l); F) \pi(q_l, x_l) + p_h^{e,j} \Psi_h(u(q_h, x_h; \theta_h); F) \pi(q_h, x_h) \\ &= \sum_{i=l,h} p_i^{e,j} \Psi_i(u(q_i, x_i; \theta_i); F) \pi(q_i, x_i)\end{aligned}$$

weighted by the probabilities that the buyer is of low or high valuation.

4.2 Seller Strategies

We so far have implicitly allowed for both pooling and separating offers; however, as our notation suggests, only separating menus are offered in equilibrium.

Corollary 4.1 (No Pooling in Equilibrium). *Equilibrium menus separate buyers of each valuation.*

Cost convexity is at the core of this result. It is the reason why a pooling contract can always be improved through a separating revision that adds quality to the high contract, at a price (above marginal costs) only high buyers are willing to pay, and reduces the quality of the low, in exchange for a discount only low buyers are interested in. We provide a simple graphical representation of this procedure in Figure 3. Beyond its economic relevance, this result is analytically convenient because it allows us to avoid both the classical threat to existence of equilibrium (a la Rothschild Stiglitz (1976)) and a more complicated diversity of offers (as in Lester et al. (2019)).

Beyond separating buyers, equilibrium menus have quite a bit of additional structure. In particular, the quality-price terms of each contract featured in a menu are closely linked to the utility they allow their targeted buyer to obtain. The forward direction is obvious since incentive constraints and individual rationality of sellers menus imply buyers will always choose their intended contract from $((q_l, x_l), (q_h, x_h))$,

$$\begin{aligned}(IC_i) : \quad & u(q_i, x_i; \theta_i) \geq u(q_{-i}, x_{-i}; \theta_i) \quad \forall i \in \{l, h\} \\ (IR_i) : \quad & u(q_i, x_i; \theta_i) \geq 0\end{aligned}$$

so the indirect utility offered to buyers of each respective valuation by a menu is simply that offered by their respective contracts,

$$\begin{aligned}u_l &= \theta_l q_l - x_l \\ u_h &= \theta_h q_h - x_h\end{aligned}$$

The backward direction follows from conditions that profit maximality imposes on optimal offers, and they imply the existence of a bijection between the indirect utilities (u_l, u_h) and quality-price terms $((q_l, x_l), (q_h, x_h))$ of contracts in any *equilibrium* menu.

Theorem 4.1 (Converting to Indirect Utilities). *Consider an equilibrium menu $((q_l, x_l), (q_h, x_h))$ with associated indirect utilities (u_l, u_h) . Qualities are then given by,*

$$q_l(u_l, u_h) = \begin{cases} \frac{u_h - u_l}{\Delta\theta} & u_h - u_l < q_l^* \Delta\theta \\ q_l^* & u_h - u_l \geq q_l^* \Delta\theta \end{cases} \quad q_h(u_l, u_h) = \begin{cases} \frac{u_h - u_l}{\Delta\theta} & u_h - u_l > q_h^* \Delta\theta \\ q_h^* & u_h - u_l \leq q_h^* \Delta\theta \end{cases} \quad (4.3)$$

and prices by $x_i = \theta_i q_i - u_i$.




Figure 3: Example with $\theta_h = 4$, $\theta_l = 2$, $\phi(q) = \frac{1}{2}x^2$. Contracts are depicted as points in the (q, x) space. Implied utilities are given by the vertical distance of a contract's y-axis coordinate to the zero utility indifference curves of the respective θ_i valuation buyers, while profits per sale are given by the vertical distance to the seller's cost function $\phi(q)$. The blue region represents the set of pooling contracts that are individually rational for buyers of either valuation, imply non-negative profits, and are not dominated by another pooling contract.

A candidate pooling offer (q_p, x_p) lies on dashed iso-utility lines for buyers with low and high valuation; we follow these to the right and left, respectively, until reaching the efficient qualities of trade with each type of buyer at prices $x_i = \theta_i q_i^* - u(q_p, x_p; \theta_i)$. This alternative separating bid $((q_l = q_l^*, x_l = x_p - \theta_l(q_p - q_l^*)), (q_h = q_h^*, x_h = x_p + \theta_h(q_h^* - q_p)))$ remains incentive compatible and strictly dominates the pooling offer: the same utility to buyers of either valuation (hence the same probability of winning) but strictly higher profits per sale.

Therefore, when an incentive constraint binds at an optimal menu, the difference in utilities $u_h - u_l$ uniquely pins down the qualities q_l, q_h offered in each contract, while the level of these utilities then uniquely determines their respective prices. When both incentive constraints are slack, however, it is optimal to offer efficient qualities to each type of buyer, which is why we will refer to these menus as *dually efficient* from now on, and the prices follow uniquely in the same fashion.

This result connecting the quality-price form of menus to their associated indirect utilities is standard in mechanism design - referred to as the *parametric-utility approach* (Rochet and Stole (2006)) - and originates from the fact that seller surplus (profits per sale) equals the gains from trade $S_i(q) = \theta_i q - \phi(q)$ net of buyer surplus,

$$\pi(q_i, x_h) = x_i - \phi(q_i) = (x_i - \theta_i q_i) + (\theta_i q_i - \phi(q_i)) = S_i(q_i) - u(q_i, x_i; \theta_i)$$

Profits then increase both from minimizing buyer surplus *and* maximizing trade efficiency. As such, consider the optimal way to offer a pair of utilities (u_l, u_h) . If it is possible to do so with a dually efficient menu that respects incentive constraints, then it is optimal to do so and unlock a larger social surplus. If it is not possible and a buyer's incentive constraint would be violated by such an offer, then a problematic IC_h (IC_l) constraint is corrected most profitably by under-providing (over-providing) quality in the low (high) contract. In these cases, a valuation θ_i buyer - whose incentive constraint binds at the optimal menu that offers (u_l, u_h) - is indifferent between their contract and the one intended for the opposite θ_{-i} valuation buyer, $u_h - u_l = q_{-i} \Delta \theta$.

Since optimal offers can be expressed in terms of indirect utilities, it is natural to recast the strategy of each type of seller $p_l^{e,j}$ as a distribution $F^{e,j}(u_l, u_h)$ over pairs of utility offers (u_l, u_h) . In this notation, the profits that a type $p_l^{e,j}$ seller expects from an offer (u_l, u_h) is,

$$\begin{aligned} \Pi_i(u_l, u_h; F) &= \Psi_i(u_i; F) \pi_i(u_l, u_h) \\ \Pi^{e,j}(u_l, u_h; F) &= \sum_{i=l,h} p_i^{e,j} \Pi_i(u_l, u_h; F) \end{aligned}$$

when the distribution of utilities is contested matches is F . Therefore, the *level* of indirect utilities (u_l, u_h) determines both the probability of winning the contested offers $F_i(u_i)$ and the part of the buyer surplus of the profits per sale $\pi(u_l, u_h) = S(u_l, u_h) - u_i$, while the *difference* in generosity towards a high and low buyer, $u_h - u_l$, determines the part of the social surplus of the profits per sale, through the efficiency of trade with them. Given the sample space $\Omega = [0, S_l^*] \times [0, S_h^*]$, Borel σ -algebra, and set of countably additive probability measures \mathcal{P} over it, this concise specification of the seller's problem allows us to introduce the intuitive equilibrium concept.

Definition 4.1 (Bayes-Nash Equilibrium). *An equilibrium is a probability measure $F \in \mathcal{P}$ over indirect utility offers (u_l, u_h) , and seller type $\{p_l^{e,j}\}_{e \in \{a,s\}, j \in \{l,h\}}$ mixed strategies $F^{e,j}$ such that,*

- *Rational Expectations:* $F(u_l, u_h) = \sum_{i \in \{l,h\}} p(\theta_i) F(u_l, u_h | \theta_i)$ where each $F(\cdot | \theta_i)$ satisfies (4.2).
- *Seller Optimality:* $\text{supp}(F^{e,j}) \subseteq \arg\max \Pi^{e,j}(u_l, u_h; F)$ for all $(e, j) \in \{a, s\} \times \{l, h\}$.

From here on, the dependence of various functions on equilibrium terms will be implied, so notation such as $(\cdot; F)$ will be suppressed and used only where it is important to avoid ambiguity.

4.2.1 Corner Cases: Monopoly and Perfect Competition

We start building intuition as in exogenous quality economies, by first analyze the equilibria of corner economies where sellers always bid alone ($\rho = 1$) or against another seller ($\rho = 0$). The screening

problem of a monopolist is well understood from Mussa Rosen (1978). The only trade-off for this seller is between the efficiency of trade with low buyers and the share of efficient social surplus S_h^* that it captures in trade with the high buyer. A type $p_l^{e,j}$ monopolist thus solves

$$\max_{u_l, u_h \geq 0} p_l^{e,j} (S_l(u_l, u_h) - u_l) + p_h^{e,j} (S_h(u_l, u_h) - u_h)$$

where incentive compatibility is implicit in the functional form of social surplus terms. We can ignore menus with (1) $u_h - u_l > q_l^* \Delta\theta$ or (2) $u_l > 0$, as they would not entail greater efficiency of trade and strictly lower profits per sale from buyers of at least one valuation than the offer $(u_l, u_h) = (0, q_l^* \Delta\theta)$. The incentive constraint of high buyers, therefore, binds at an optimal menu, and the optimal quality in the low contract, for a type $p_l^{e,j}$ seller with strictly convex costs, is determined by the marginal condition,

$$\underbrace{p_l^{e,j} \cdot \frac{\theta_l - \phi'(q_l(0, u_h^{e,j}))}{\Delta\theta}}_{\text{efficiency gains}} - \underbrace{p_h^{e,j}}_{\text{rent losses}} = 0 \quad (4.4)$$

The left-hand term is the marginal gain in efficiency made possible by increasing the indirect utility offered to high buyers. This relaxes their incentive constraint and allows the seller to provide a more efficient quality to low buyers, increasing the social surplus from trade with them and thus increasing the profitability of each low sale. The right-hand term represents the rents surrendered by making a more generous offer to high buyers - featuring the same quality at a strictly lower price. Monopolists weigh these marginal effects by their type, the conditional probability $p_l^{e,j}$ of being in a low match.

Lemma 4.1 (Monopoly). *If $\rho = 1$, a seller of type $p_l^{e,j}$ with strictly convex costs offers a menu $(u_l^{e,j}, u_h^{e,j})$ with $u_l^{e,j} = 0$ and $u_h^{e,j}$ satisfying Equation (4.4).*

Consider two monopolists of types $p_l^{e,j} < p_l^{e',j}$ and note that (4.4) necessarily implies the latter is more generous, $u_h^{e,j} < u_h^{e',j}$. This means that monopolists' bids are monotonically ranked (*ordered*) by type, a property that we saw in exogenous quality economies where some rationed offers were extended, and that will also generalize to endogenous quality economies featuring interior levels of crowding ($\rho \in (0, 1)$). We can also easily characterize the impact of expertise on all agents in the monopoly setting,

- **Buyers:** High buyers prefer to face a seller with less expertise, whereas low buyers are indifferent.
- **Sellers:** Expertise strictly increases ex-ante profits.
- **Social Planner:** Expertise strictly improves the efficiency of trade: high buyers always receive efficient offers, but that of the low is increasing in expertise.

These relations extend to interior levels of competition with an essential caveat: competitive pressure allows low buyers to benefit from expertise and partake in some of the additional social surplus it generates.

The other corner where sellers are guaranteed to bid against another seller in every match ($\rho = 0$) is also well understood, as each match again becomes a setting of Bertrand price competition. Therefore, sellers make dually efficient offers, and buyers capture the entire social surplus.

Theorem 4.2 (Perfect Competition). *If $\rho = 0$, the unique Bayes-Nash equilibrium involves almost every seller offering $((q_l, x_l), (q_h, x_h)) = ((q_l^*, \phi(q_l^*)), (q_h^*, \phi(q_h^*)))$ with probability 1.*

On one hand, this implies the negative point that perfect competition is the wrong setting for the study of predictive skill about buyer's willingness to pay for vertically differentiated goods, mainly

because it is irrelevant. On the other hand, it implies the positive point that expertise is not valuable if sellers cannot use it to improve the profitability of their offers and that the equilibrium structure has a large impact on this. In interior economies, crowding will also promote efficiency and consumer surplus, but it will differentially shape the returns to expertise of amateurs and sharks because the latter's offers will often be dually efficient and thus unimprovable at the margin by more precise demand information. Lastly, we highlight that a strategic dependence between sellers' bids and those of peers requires some degree of crowding, so interior levels, where crowding is not so strong as to render expertise moot, will allow us to study the externalities of expertise on other sellers. We will find that sellers become more profitable through their own expertise and even sometimes that of their peers.

4.3 Equilibrium Structure

In this section, we discuss some general properties satisfied by the menus of a candidate equilibrium. Each has intuitive appeal, either from an economic standpoint or because of the mathematical tractability that they impart. By leveraging these in conjunction with the optimality conditions of sellers' problems, we can solve for this equilibrium analytically and obtain a complete characterization. In the Appendix, we show that these properties hold in any equilibrium where at least some offer ratios low buyers and that, in fact, these economies have a unique equilibrium - the conjectured one.

4.3.1 Claims about Equilibrium Structure

There are three main areas that benefit from additional structure: (1) the distributions over indirect utilities offered by sellers in low and high matches, F_i , (2) the relationship of incentive compatibility constraints to the generosity (indirect utility) of menus, and (3) the relationship between generosity of a menu with the type $p_i^{e,j}$ of the seller who offers it.

The equilibrium distributions of indirect utilities F_i offered in each match are weighted averages of each seller type's mixed strategy, and it inherits their properties.

Claim 4.1. *The equilibrium marginal distributions over indirect utilities F_i for $i \in \{l, h\}$,*

1. *are atomless.*
2. *have a connected support of low utility offerings $\Upsilon_l = [\underline{u}_l, \bar{u}_l]$.*
3. *have a support of high utility offerings $\Upsilon_h = \cup_{e \in \{a, s\}, j \in \{l, h\}} [\underline{u}_h^{e,j}, \bar{u}_h^{e,j}]$, made up of the contiguous bids by type $p_i^{e,j}$ sellers.*
4. *are continuously differentiable with densities f_i in the interior and one-sided derivatives at the boundaries.*

By avoiding atoms and gaps in the supports of each seller type's mixtures, erratic behavior is curtailed. It is never optimal to bunch up and compete at a single point in the space of utility offers, as one might see in the corner cases of monopoly and Bertrand, nor do we see discontinuous jumps in generosity among sellers of the same type. While continuous densities allow us to consider marginal incentives, which helps convey economic intuition and solve for the equilibrium as the solution of a standard differential system.

As in the closely related work of Lester et al. (2019) and Garret et al. (2019), a property called *ordering*, which refers to an equilibrium where ranking menus by the utility offered to high or low

buyers is identical, lends a great degree of tractability. This property turns out to be necessary in any equilibrium where low buyers' incentive constraint is slack in every offered menu, and a sufficient condition for this is the assumption that sellers cannot profitably trade the efficient high quality with low buyers.

Assumption 4.2 (No IC_l Cost Condition). *The piecewise linear cost function $\phi(\cdot)$ is such that,*

$$\theta_l q_h^* \leq \kappa_m(q_h^* - q_l^*) + \kappa_l q_l^* \quad (4.5)$$

Claim 4.2. *Low buyer's incentive constraint never binds at equilibrium menus.*

Eliminating the possibility of a binding low buyer constraint preserves efficiency at the top, so high buyers always obtain their efficient quality (q_h^*), but low buyers are rationed ($q_l < q_l^*$) whenever the offer intended for them is part of a menu in which high buyers' incentive constraint binds, as in the monopoly setting.

We will focus on equilibria in which offers are (a) ordered by their generosity, but also (b) monotone in the seller's type, meaning that generosity increases with the seller's posterior probability of being matched with a low buyer $p_l^{e,j}$, as in the ordered equilibria of economies where quality is exogenous.

Claim 4.3. *Given two equilibrium menus (u_l, u_h) and $(\tilde{u}_l, \tilde{u}_h)$ offered by sellers of respective types p_l and \tilde{p}_l with $u_i > \tilde{u}_i$ for some $i \in \{l, h\}$,*

1. *It is also true that $u_{-i} \geq \tilde{u}_{-i}$ and strictly so if both menus are offered by sellers of the same type.*
2. *The gap between low and high utilities $u_h - u_l$ increases strictly with generosity.*
3. *Each support $\Upsilon_i = [\underline{u}_i, \bar{u}_i]$ is such that $\bar{u}_i \leq S_i^*$. Further, there exists a $u_i^{de} \in [\underline{u}_i, \bar{u}_i]$ such that all $u_i < u_i^{de}$ are in IC_h binding menus and all $u_i \geq u_i^{de}$ are dually efficient.*

Due to Theorem 4.1, the efficiency of any menu is directly pinned down by the difference in its respective utility offers $u_h - u_l$, and since low buyers' incentive constraint does not bind in the menus of these equilibria, a larger utility difference only improves efficiency - weakly raising the quality of low trade towards the optimal, q_l^* . Orderedness in generosity then follows from the complementarity between relaxing the incentive compatibility constraint of a high buyer (through a more generous u_h term) - which permits more profitable sales to low buyers - and increasing low sales (through a more generous u_l term). In equilibrium, the first effect dominates - sellers face stronger incentives to increase the generosity of their high vs low offer - so that the difference in utilities $u_h - u_l$ and, consequently, efficiency is weakly increasing in generosity. This upward efficiency progression then also generates a natural grouping of menus, with the least generous menus also being the least efficient (constrained by high buyer incentive compatibility), but above a generosity threshold, becoming dually efficient (unconstrained by either buyer's incentive compatibility). The monotonic relation between a seller's type and generosity, on the other hand, is due to the link between generosity and profitability of equilibrium contracts in a match with each kind of buyer.

Claim 4.4. *Given two equilibrium menus (u_l, u_h) and $(\tilde{u}_l, \tilde{u}_h)$ offered by sellers of respective types p_l and \tilde{p}_l with $u_i > \tilde{u}_i$ for some $i \in \{l, h\}$,*

1. *Profitability conditionally on matching with a buyer of low (high) valuation is increasing (decreasing) in generosity,*

$$\Pi_l(u_l, u_h) > \Pi_l(\tilde{u}_l, \tilde{u}_h) \text{ and } \Pi_h(u_l, u_h) > \Pi_h(\tilde{u}_l, \tilde{u}_h)$$

2. *Generosity is increasing in the seller's conviction of facing a low valuation buyer, $p_l > \tilde{p}_l$, and strictly so if high buyers' incentive constraint binds at (u_l, u_h) .*

Equilibrium profits conditionally on matching with a low (high) buyer provide the incentives that order the offers of sellers with different interim assessments (types). In particular, the profitability of menus in low (high) matches will increase (decrease) with their generosity. The origin of this relation is clear when we consider the first order condition satisfied by the bids of sellers offering menus at which high buyers' incentive constraint binds,

$$0 = p_h^{e,j} \frac{\partial \Pi_h}{\partial u_h}(u_h, u_l) + p_l^{e,j} \underbrace{\frac{\partial \Pi_l}{\partial u_h}(u_h, u_l)}_{\geq 0}$$

$$0 = p_l^{e,j} \frac{\partial \Pi_l}{\partial u_l}(u_h, u_l)$$

When high buyers' incentive constraint binds at a menu, high generosity relaxes this constraint and makes low trade more profitable, so the term $\frac{\partial \Pi_l}{\partial u_h}(u_h, u_l)$ is strictly positive. As such, the profits from high sales of constrained menus (Π_h) are locally decreasing in generosity towards high buyers (u_h) but invariant in that towards the low (u_l), while their profits from low sales (Π_l) are increasing generosity towards high buyers and at a local maximum with respect to low buyer generosity. When we consider the set of constrained menus, the least generous ones are offered by the sellers who are more relatively more concerned about the profitability in high matches, mainly those sellers with the lowest type $p_l^{e,j}$. The efficiency-generosity relation places these constrained menus below any dually efficient ones, so sellers who are sufficiently concerned about low match profits offer the latter (dually efficient) menus. However, when menus are dual efficient, incentive constraints are slack, and changing the generosity of the contract offered to a buyer of either valuation does not impact the profitability of the paired contract, offered to a buyer of the other valuation (by Theorem 4.1). The term $\frac{\partial \Pi_l}{\partial u_h}(u_h, u_l)$ is, therefore, equal to zero at dually efficient offers, so these satisfy the pair of equations,

$$0 = \frac{\partial \Pi_h}{\partial u_h}(u_h, u_l)$$

$$0 = \frac{\partial \Pi_l}{\partial u_l}(u_h, u_l)$$

which implies that every dually efficient menu has identical profits in a match with a θ_i buyer, and so that all sellers - irrespective of their type - are indifferent between these menus. The most we can necessarily say about the relation between a seller's type and dually efficient menus is that the type must be high enough for the seller to offer any dually efficient menus. Any additional relation between sellers' type and the generosity of dually efficient menus is not just unnecessary and unintuitive but unappealing because it would allow expertise to shape the share of surplus that buyers of each valuation obtain - in a sense, generating effects of expertise where there don't have to be any. This is an important difference with economies where quality is exogenous, since the generosity (as reflected by their price) of offers in these, including of efficient ones that allow trade with low buyers, is strictly monotone in the seller's type when the equilibrium features some constrained (inefficient) offers.

We close by connecting these points about profitability, generosity, and a seller's type to the benefits of expertise. In our setting, with two information structures (levels of expertise), each of which features a random variable (signal) that can take two possible values (low or high), there are four types of sellers at the interim stage, corresponding to each expertise and signal combination. Naturally, sellers with the highest expertise have the greatest conviction about their signals, and

receiving a high (low) signal leads any seller to revise their prior down (up), so types are ranked as per,

$$p_l^{s,h} < p_l^{a,h} < p_l^{a,l} < p_l^{s,l}$$

So that expertise has any value, consider an equilibrium where some inefficient menus are offered. The least generous menus, towards a buyer of either valuation, are offered by sharks who observe high signals, while the most generous are offered by sharks who observe low signals. Amateur offers have intermediate generosity and only overlap with those of sharks if they are dually efficient. As such, in low (high) matches, a greater share of sharks (than amateurs) observe signals that orient their offers to the top (bottom) of the generosity distribution, where the most profitable low (high) match menus are found. This makes sharks more profitable in the average match than amateurs - the benefit of expertise.

4.4 Comparative Statics

We will perform an analogous comparative static analysis to Section 3.4, relating points of commonality and departure, thereby tracing out where product design makes its mark. The aggregate statistics for welfare and trade efficiency are similar after adjusting for the fact that we express strategies in the space of indirect utility offers in these settings with endogenous quality choices. Recycling notation, the welfare of each type of buyer follows as,

$$\mathcal{W}_b(\theta) = \rho E[u_i] + (1 - \rho) E[\max(u_i, \tilde{u}_i)] \quad (4.6)$$

where u_i are iid draws from the conditional equilibrium marginals $F_i(u_i|\theta_i)$. Seller profits are given by,

$$\Pi^e = E \left[\underbrace{\mathcal{M}}_{\text{number of matches}} \underbrace{\Psi_i(u_i)}_{\text{probability of selling}} \underbrace{S_i(u_l, u_h) - u_i}_{\text{profits-per-sale}} \middle| e \right] \quad (4.7)$$

where we also average over the probabilities that the seller observes each signal $j \in \{l, h\}$ in a match with a buyer of each valuation $\theta_i \in \{\theta_l, \theta_h\}$, and that it offers the menus $(u_l, u_h) \sim F^{e,j}$ when they of either type $p_l^{e,j}$. Lastly, in the case of efficiency, high buyers always receive efficient offers, but the low may get rationed, so we measure the efficiency of trade by the average quality offered to the latter,

$$\rho E[q_l|\theta_l] + (1 - \rho) E[\max(q_l, q'_l)|\theta_l] \quad (4.8)$$

where we rely on the equilibrium property that generosity and efficiency are positively correlated, so buyers always select an offer with the greatest efficiency available.

Flexible production introduces changes that are profound, yet easy to overlook because they can be drowned out by the overarching point that the qualitative effects of crowding and expertise on efficiency and welfare aggregates do not change - directionally similar, under similar environments. However, closer inspection reveals the fundamental differences that information-driven production decisions introduce to welfare analysis. For one, it changes where we should look to track efficiency. Whereas pricing uniquely impacted trade efficiency at the extensive margin - determining whether a buyer trades or not by finding any acceptable prices - when production was inflexible (exogenous), removing this barrier allows information to shape the level of trade at the intensive margin - determining the quality of trade a buyer obtains. Stronger still, the optimality of screening menus results in almost every offer allowing *some* low trade, and so that *all* action is at the intensive margin - a theoretical insight that informs empirical analyses. Further, expertise heterogeneity becomes a central issue, as

it is potentially intrinsic to these settings, and it is only by including it in our model that we find the novel reason why. In particular, we find that competition affects amateurs and sharks differently, making the former's (latter's) offers less (more) sensitive to additional precision; therefore, competition neutralizes amateur expertise - making it less hurtful for other sellers and discouraging its acquisition - and amplifies shark expertise - making it more beneficial for all sellers and encouraging its acquisition.

4.4.1 Crowding

The aggregate effects of crowding are similar when we allow information to also orient production - increasing trade efficiency, buyers' share of surplus, and decreasing sellers' profits - but the manner in which these arise differs. Mechanically, this set of results is linked by the fact that crowding increases the sales gains from generosity, so sellers offer more generous menus, and complementarity between the utility offers to each type of buyer both links their joint progression and that of the low-quality good. Buyers almost always trade in these settings, so efficiency increases because of action at the intensive margin (volume of trade). This is a significant difference from the usual case in the literature, and also in our model when quality is exogenous, where pricing uniquely shapes efficiency - either pricing buyers out or not - and action is at the extensive margin (probability of trade).




Figure 4: The common parameters are $[\theta_l, \theta_h, p_l] = [1, 3, 0.5]$ for buyers, $[\kappa_l, \kappa_m, q_l^*, q_h^*] = [0.5, 1, 2, 5]$ for sellers' cost functions, $[\alpha_a, \alpha_s] = [0.55, 0.95]$ for sellers' information technologies, and $p(a) = 0.5$ for proportion of amateurs.

4.4.2 Expertise

The aggregate qualitative effects of expertise are also similar. In fact, the comparison of figures under exogenous, Table 2 - Table 4, and endogenous quality, Table 5-Table 7, reveals a striking resemblance down to the very marginal effects of expertise (curvature). But, upon closer inspection, we do notice some significant differences.

When it comes to buyers, the principal difference of endogenizing quality - beyond a direct change in the variety of goods - is that when sellers cannot screen, they agree - irrespective of their type - on the optimal inefficient (separating) offers, whereas when sellers can screen, they agree - irrespective of their type - on the optimal (dually) efficient offers. Under sufficient crowding, amateurs offer efficient offers in both types of economies; so, in the former, expertise still influences their preferred bids, but in the latter, it does not. This is why in the lower left-hand figure of Table 2 their expertise has aggregate effects, but in Table 5 it does not.

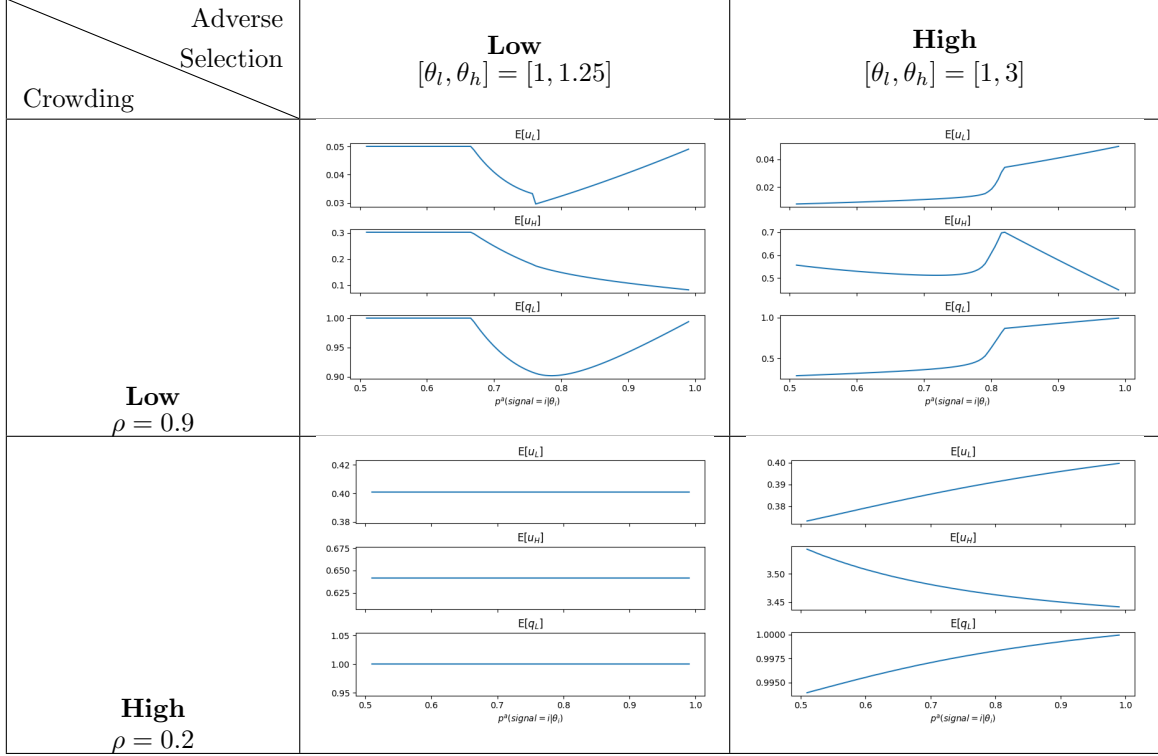


Table 5: Utility and Efficiency Effects of Precision: Almost all sellers are amateurs (90%), and we increase the precision of their signals along the respective x-axes. Each figure is representative of the effect that precision has on buyer surplus and trade at different levels of crowding and adverse selection. The common parameters are $p(\theta_l) = 0.5$ for the mass of low buyers, $[\kappa_l, \kappa_m, q_l^*, q_h^*] = [0.5, 1.2, 2, 1, 2]$ for sellers' cost functions, and $\alpha_s = 0.999$ for shark precision.

When it comes to sellers, the flip side of competition making amateur offers less sensitive to their information is that amateur expertise becomes less threatening to sharks. Concretely, we see this in the lower rows of Table 3 and Table 6, where amateur expertise is less damaging to shark profits. Whereas the flip side of sellers not agreeing on the inefficient offers under endogenous quality choices is that shark expertise more strongly relaxes competition during high matches, because it doesn't just change the mass of them that prefer to extend inefficient offers (as is also the case under exogenous quality), but also the terms of those inefficient offers - becoming stringier as their confidence in a high match increases. As a result, if we compare Table 4 and Table 7, it is evident that shark expertise is more beneficial for amateurs when it can also inform production decisions.




Table 6: Profit Effects of Amateur Precision: A supermajority of sellers are amateurs (90%) and we increase the precision of their signals along the respective x-axes. Each figure is representative of the effect that precision has on their and sharks' profits in low matches, high matches, and the average match. The common parameters are $p(\theta_l) = 0.5$ for the mass of low buyers, $[\kappa_l, \kappa_m, q_l^*, q_h^*] = [0.5, 1.2, 2, 1, 2]$ for sellers' cost functions, and $\alpha_s = 0.999$ for shark precision.




Table 7: Profit Effects of Shark Precision: A supermajority of sellers are sharks (90%) and we increase the precision of their signals along the respective x-axes. Each figure is representative of the effect that precision has on their and amateurs’ profits in low matches, high matches, and the average match. The common parameters are $p(\theta_l) = 0.5$ for the mass of low buyers, $[\kappa_l, \kappa_m, q_l^*, q_h^*] = [0.5, 1.2, 2, 1, 2]$ for sellers’ cost functions, and $\alpha_a = 0.55$ for amateur precision.

5 Conclusion

A race in analytics is currently taking place, where firms’ competitive performance is intrinsically linked to their adeptness in sourcing and utilizing data. Although it may be too early to tell, an emerging literature¹⁷ argues that many of the technologies being deployed - machine learning, big data, data mining, natural language processing, etc. - are indeed general-purpose technologies (GPT) i.e. “widely used, capable of ongoing technical improvement, and enabling innovation in application sectors” (Bresnahan et al. (1995), Bresnahan (2010)), and so in need of additional research attention given the significant ramifications of their deployment in markets.

Motivated by the size of the retail sector and its intense reliance on data to optimize product and pricing decisions, we answer a “crucial” question posed by Bergemann and Bonatti (2019) in their review of markets for information: “what are the implications of acquiring an advantage in a downstream market by means of better data (e.g. improvements in the predictive power of an algorithm)?” The problem is addressed within a model that features sellers with predictive skill, referred to as *expertise*, and a strategic role for it, as competition is imperfect. Crucially, we allow ex-ante heterogeneity in expertise to match the empirical documentation of broad, endemic differences amongst firms in this capability. We find expertise to be generically efficiency enhancing though redistributive. On the demand side, it tends to benefit low valuation buyers but hurt high valuation ones. On the supply side, unsurprisingly, expert firms obtain larger profits, but more interestingly, their expertise can also benefit competitors through externalities in the efficiency of trade. This

¹⁷Brynjolfsson, Rock and Syverson (2019), Cockburn et al. (2019), Trajtenberg (2019), Goldfarb (2019)

suggests an ambiguous distributional impact, in contrast to the unequivocal harm that many critics of big data and analytics anticipate.

Several compromises made in our modeling have left open avenues for future research with connections to concurrent work. For one, expertise is modeled in a very reduced form, which ignores crucial factors like the fundamental feedback loop between data and the skills it enables as well as the strategic choices in expertise acquisition. These choices handcuff our model to be static and the composition of the supply side to be exogenous in its most natural interpretation, ultimately limiting us to suggestive conjectures about the nature of competition and its dynamics. A formal characterization of these would be very helpful for understanding whether heterogeneity is a short-run phenomenon or a fundamental feature of some steady-states. Lastly, the naivete of buyers with respect to sellers' information could also benefit from relaxation given its importance to outcomes, even within the present setting, as well as its role in a host of results emerging from the literature on consumer privacy. Much like these agents, we are using a model to understand a complicated problem, however, structural concerns have led us to sacrifice additional complexity for the sake of intelligibility. Our hope is that this work contributes to an ongoing discussion that advances our joint expertise.

Appendix A Exogenous Quality Setting: Seller Strategies

A.1 Economies with Separating and Pooling Offers

Starting with the strategy of seller types that offer the (weakly) highest prices (sharks who observe high signals), the supremum of their distribution's support is as specified by Proposition 3.7. If their highest price is θ_l , we focus on the symmetric equilibrium where all sellers have the same mixture to avoid introducing an arbitrary role for expertise. When every type of seller has the same strategy however, the distribution of competitor offers in both low and high matches is precisely equal to it, so a single equation at every price offer pins down this distribution $F(x)$; mainly, that which guarantees that every offer yields the same profits as the highest one,

$$\rho(\theta - \kappa) = (\rho + (1 - \rho)(1 - F(x)))(x - \kappa) \quad (\text{A.1})$$

Proposition A.1 (All Pooling Symmetric Equilibrium). *If the highest equilibrium is pooling, then the distribution of competition offers in low and high matches, as well as, each type of seller's strategy is given by the distribution,*

$$F(x) = 1 - \frac{\rho}{1 - \rho}(\theta - x) \quad (\text{A.2})$$

However, when the highest equilibrium price is separating, we study the unique ordered equilibrium, without loss of generality, and since seller strategies are ordered in these equilibria, we solve for each of these differing strategies sequentially, starting with that of sharks who observe high signals. In the interest of clarity, we will first derive strategies assuming that no seller type offers both separating and pooling prices; then we will show how to modify the solution when the assumption is violated. Since all prices of sharks who observe high signals are separating in these economies, their lowest price offer is also separating and given by,

$$\begin{aligned} \rho(\theta_h - \kappa) &= (\rho + (1 - \rho)\mu(s)P^s(h|\theta_h))(\underline{x}^{s,h} - \kappa) \\ \implies \underline{x}^{s,h} &= \kappa + \frac{\rho}{\rho + (1 - \rho)\mu(s)P^s(h|\theta_h)}(\theta_h - \kappa) \end{aligned} \quad (\text{A.3})$$

where $\mu(s)P^s(h|\theta_h)$ is the mass of peer sharks who observe h that a seller expects to face in competitive matches if the buyer is of high valuation - it would always lose against sellers of higher type, as they offer lower prices in the ordered equilibrium. The allocation of mass between the upper and lower bound is then uniquely determined by (3.4) and (3.12),

$$F^{s,h}(x) = \begin{cases} 1 & \forall x \in (\theta_h, \infty) \\ 1 - \frac{\rho}{(1 - \rho)P(p_l^{s,j}|\theta_h)} \frac{\theta_h - x}{x - \kappa} & \forall x \in [\underline{x}^{s,h}, \theta_h] \\ 0 & \forall x \in (-\infty, \underline{x}^{s,h}) \end{cases} \quad (\text{A.4})$$

$$P(p_l^{s,j}|\theta_h) = \mu(s)P^s(h|\theta_h) \quad (\text{A.5})$$

By induction, the strategy of any other type of seller $p_l^{e,j}$ who only makes separating offers,

$$F^{e,j}(x) = \begin{cases} 1 & \forall x \in (\bar{x}^{e,j}, \infty) \\ \left(1 - \frac{P^e(\bar{x}^{e,j} \text{ sale}|\theta_h)}{(1-\rho)P^e(p_l^{e,j}|\theta_h)} \frac{\bar{x}^{e,j}-x}{x-\kappa}\right) & \forall x \in [\underline{x}^{e,j}, \bar{x}^{e,j}] \\ 0 & \forall x \in (-\infty, \underline{x}^{e,j}) \end{cases} \quad (\text{A.6})$$

$$P(p_l^{e,j}|\theta_h) = \mu(e)P^e(j|\theta_h) \quad (\text{A.7})$$

$$P(x \text{ sale}|\theta_h) = \rho + (1-\rho) \sum_{\substack{e' \in \{a,s\} \\ j' \in \{l,h\}}} \mu(e')P^{e'}(j'|\theta_h)(1 - F^{e',j'}(x)) \mathbb{1}(p_l^{e',j'} \leq p_l^{e,j}) \quad (\text{A.8})$$

$$\bar{x}^{e,j} = \underline{x}^{e,j,-} \quad (\text{A.9})$$

$$\underline{x}^{e,j} = \kappa + \frac{P^e(\bar{x}^{e,j} \text{ sale}|j, \theta_h)}{P^e(\underline{x}^{e,j} \text{ sale}|j, \theta_h)}(\bar{x}^{e,j} - \kappa) \quad (\text{A.10})$$

where $p_l^{e,j,-}$ is the closest type from below. The terms $(1 - F^{e',j'}(x)) \mathbb{1}(p_l^{e',j'} \leq p_l^{e,j})$ in (A.8) are uniquely determined by the inductive assumption, if they are the strategies of lower type sellers ($p_l^{e',j'} < p_l^{e,j}$), or otherwise equal to zero, if are the strategies of higher type sellers ($p_l^{e',j'} > p_l^{e,j}$). As such, the strategy of the $p_l^{e,j}$ type seller, $F^{e,j}(x)$, is also uniquely determined.

Continuing with sellers who offer pooling prices. The highest such price is $\bar{x}^{e,j} = \theta_l$ (Proposition 3.8), and the form of upper and lower bounds as well as strategies of the types of sellers who make pooling offers are analogous to those of types who make separating offers, after adjusting the type's expected sales to include its posterior probability of being in a low match,

$$F^{e,j}(x) = \begin{cases} 1 & \forall x \in (\bar{x}^{e,j}, \infty) \\ 1 - \frac{P^e(\bar{x}^{e,j} \text{ sale}|j)}{(1-\rho)P^e(p_l^{e,j}|\theta_h)} \frac{\bar{x}^{e,j}-x}{x-\kappa} & \forall x \in [\underline{x}^{e,j}, \bar{x}^{e,j}] \\ 0 & \forall x \in (-\infty, \underline{x}^{e,j}) \end{cases} \quad (\text{A.11})$$

$$P^e(p_l^{e,j}|\theta_h) = p_l^{e,j} \mu(e)P^e(j|\theta_l) + p_h^{e,j} \mu(e)P^e(j|\theta_h) \quad (\text{A.12})$$

$$P^e(x \text{ sale}|j) = \rho + (1-\rho) \sum_{\substack{j' \in \{l,h\}, \\ e' \in \{a,s\}}} \mu(e') \left(p_l^{e,j} P^{e'}(j'|\theta_l) + p_h^{e,j} P^{e'}(j'|\theta_h) \right) (1 - F^{e',j'}(x)) \mathbb{1}(p_l^{e',j'} \leq p_l^{e,j}) \quad (\text{A.13})$$

$$\bar{x}^{e,j} = \underline{x}^{e,j,-} \quad (\text{A.14})$$

$$\underline{x}^{e,j} = \kappa + \frac{P^e(\bar{x}^{e,j} \text{ sale}|j)}{P^e(\underline{x}^{e,j} \text{ sale}|j)}(\bar{x}^{e,j} - \kappa) \quad (\text{A.15})$$

It is understood that $\underline{x}^{e,j,-} = \theta_l$ for the type of seller that offers the highest pooling price.

A.2 Seller Strategies with Disjoint Support

We can verify if a seller type who offers separating prices also offers any pooling ones, by checking if offering the highest possible pooling price (θ_l), and beating all sellers of equal or lower type in contested matches, it could make strictly greater profits than those of monopoly offers,

$$p_j^{e,h} \rho(\theta_h - \kappa) < \left(\rho + (1-\rho) \sum_{p_l^{e',j'} \leq p^{e,j}} (p_l^{e,j} P(p_l^{e',j'}|\theta_l) + p_h^{e,j} P(p_l^{e',j'}|\theta_h)) \right) (\theta_l - \kappa)$$

$$\frac{\rho}{(1-\rho) \sum_{p_l^{e',j'} \leq p^{e,j}} (p_l^{e,j} P(p_l^{e',j'}|\theta_l) + p_h^{e,j} P(p_l^{e',j'}|\theta_h))} < \frac{\theta_l - \kappa}{p_j^{e,h}(\theta_h - \kappa) - (\theta_l - \kappa)} \quad (\text{Discontinuity Condition})$$

When this condition holds, this type's separating distribution continues as specified by (A.6), up to its lowest separating offer $\hat{x}_l^{e,j}$ given by,

$$\begin{aligned} & p_h^{e,j} \left(\rho + (1-\rho) \sum_{p_l^{e',j'} < p^{e,j}} P(p_l^{e',j'} | \theta_h) \right) (\bar{x}^{e,j} - \kappa) \\ &= \left(\rho + (1-\rho) \sum_{p_l^{e',j'} \leq p^{e,j}} (p_l^{e,j} P(p_l^{e',j'} | \theta_l) + p_h^{e,j} P(p_l^{e',j'} | \theta_h)) (1 - F^{e',j'}(\hat{x}^{e,j})) \right) (\theta_l - \kappa) \end{aligned} \quad (\text{A.16})$$

since no other type offers prices in the gap $[\theta_l, \hat{x}_l^{e,j}]$, nor are there any point masses in equilibrium mixtures, whereby $F^{e,j}(\theta_l) = F^{e,j}(\hat{x}^{s,h})$. As such, we can solve for $\hat{x}^{e,j}$ explicitly by substituting the form of the distribution over separating offers (A.6),

$$\begin{aligned} & \hat{x}^{e,j} \\ &= \frac{\bar{x}^{e,j}}{1 + \frac{p_h^{e,j} \left(\rho + (1-\rho) \sum_{p_l^{e',j'} < p^{e,j}} P(p_l^{e',j'} | \theta_h) \right) \frac{\bar{x}^{e,j}}{\theta_l} - \left(\rho + (1-\rho) \sum_{p_l^{e',j'} < p^{e,j}} (p_l^{e,j} P(p_l^{e',j'} | \theta_l) + p_h^{e,j} P(p_l^{e',j'} | \theta_h)) \right)}{(p_l^{e,j} P(p_l^{e,j} | \theta_l) + p_h^{e,j} P(p_l^{e,j} | \theta_h)) \frac{P(\bar{x}^{e,j} | \text{sales} | j)}{P(p_l^{e,j} | \theta_h)}}}} \end{aligned} \quad (\text{A.17})$$

The mixture continues downwards over pooling prices in the standard manner, as per (A.11), until the lowest overall price,

$$\begin{aligned} & p_h^{e,j} \left(\rho + (1-\rho) \sum_{p_l^{e',j'} < p^{e,j}} P(p_l^{e',j'} | \theta_h) \right) \bar{x}^{e,j} \\ &= \left(\rho + (1-\rho) \sum_{p_l^{e',j'} \leq p^{e,j}} (p_l^{e,j} P(p_l^{e',j'} | \theta_l) + p_h^{e,j} P(p_l^{e',j'} | \theta_h)) \right) \underline{x}^{e,j} \\ &\implies \underline{x}^{s,h} = \frac{p_h^{e,j} \left(\rho + (1-\rho) \sum_{p_l^{e',j'} < p^{e,j}} P(p_l^{e',j'} | \theta_h) \right)}{\rho + (1-\rho) \sum_{p_l^{e',j'} \leq p^{e,j}} (p_l^{e,j} P(p_l^{e',j'} | \theta_l) + p_h^{e,j} P(p_l^{e',j'} | \theta_h))} \end{aligned} \quad (\text{A.18})$$

Appendix B Endogenous Quality Setting: Equilibrium Properties

In this section, we will discuss the technical aspects of the candidate equilibrium introduced in Section 4.3 and explain why many of its distinguishing properties hold in all equilibria. An even stronger result will follow that the ordered symmetric equilibrium is unique, which allows us to affirm the genericity of the comparative static analysis in Section 4.4.

B.1 Equilibrium Distributions and Orderedness

We begin by recalling the concept of an ordered equilibrium, which connects the level of indirect utility offered to low and high buyers in a menu. Put simply, in an ordered equilibrium, sellers who offer more indirect utility in the contract intended for a high valuation buyer must do the same in the contract intended for a low valuation one.

Definition B.1 (Orderedness). *An equilibrium is said to be weakly-ordered if, for any two equilibrium menus (u_l, u_h) and (u'_l, u'_h) ,*

$$(u_h - u'_h)(u_l - u'_l) \geq 0$$

When the inequality holds strictly in almost every¹⁸ comparison, we refer to the equilibrium as ordered.

These two properties have also been referred to as rank preserving and strictly rank preserving in related work (including Lester et al. (2019)). Like them, we find that orderedness is necessary holds whenever a buyer's incentive compatibility constraint binds at one of the menus that is being compared.

Lemma B.1 (Ordered Equilibrium). *Almost every equilibrium menu (u_l, u_h) featuring a binding incentive compatibility constraint is ordered when compared to another equilibrium menu (u'_l, u'_h) . That is,*

$$(u_h - u'_h)(u_l - u'_l) > 0$$

The economic rationale underlying this complementarity in indirect utilities is familiar from Garrett et al. (2019) and intuitively explained in Section 4.3. Consider a seller who increases the indirect utility it offers to high buyers u_h . If their incentive constraint binds, this relaxes it and allows for an increase in the quality provided to buyers of low valuation in the paired contract (q_l, x_l) , as per $q_l = \frac{u_h - u_l}{\Delta\theta}$. The seller can then offer low buyers the same utility u_l , while obtaining strictly larger profits in each sale to them; however, when low sales are more profitable, the seller also bids more aggressively for them, and this is done by making them more appealing through a low utility increase ($\uparrow u_l$). The channel that creates this complementarity is, therefore, the connection between the efficiency of the contracts and the utility they offer to both buyers. Since this link is missing among dually efficient offers, at which incentive compatibility constraints are slack, they do not need to be ordered.

Orderedness simplifies the equilibrium structure substantially, and heterogeneity in sellers' interim beliefs does not alter the fundamental complementarity between providing additional utility to low and high buyers, but rather how interested sellers are in forfeiting high sale profitability for low one. Therefore, the heterogeneity of the posteriors *moderates* the joint progression of u_l and u_h , but does not change the correlation between these.

Theorem B.1 (Type Monotonicity). *Let (u_l, u_h) and (u'_l, u'_h) be two equilibrium menus sharing a common binding incentive compatibility constraint with $u_i < u'_i$. Then,*

1. *High sale profits per-match are decreasing in generosity, $\Pi_h(u_l, u_h) > \Pi_h(u'_l, u'_h)$, while low ones increase, $(u_l, u_h) < \Pi_l(u'_l, u'^j_h)$.*
2. *If the menus are offered by sellers of respective types $p_l \neq p'_l$, then $p_l < p'_l$*

Consider the menu $(u_l(p), u_h(p))$ occupying the p^{th} generosity-percentile in the equilibrium distribution. Then, $u_h(p)$ increases enough for profits from high sales $\Psi_h(p)(S_h^* - u_h(p))$ to decrease, but $u_l(p)$ grows passively enough to not undo the additional profitability of profits from low sales. The relationship between low/high trade profits per-match and generosity creates an equilibrium structure where seller types comparatively more interested in profits from high sales make offers that are less generous towards buyers with either valuation than those made by seller types comparatively more interested in profits from low sales. In particular, sellers relatively more convinced that they face high buyers will aim to depress bids as much as possible so as to extract these buyers' (information) rents, whereas sellers who are relatively more convinced that they face low buyers give greater consideration

¹⁸Up to a measure zero set of menus.

to capturing profitable trade with them and cede additional rents to both low and high buyers to do so.

Efficiency gains thus far have been described as taking place within low trade, implicitly treating the incentive constraint of high buyers as the only relevant one. In fact, this is a necessary property of equilibria in any economy where sellers' costs satisfy Assumption 4.2. Furthermore, in these, offers are grouped in two sets of menus. The most rationed menus, at which high buyers' incentive constraints binds, are also the least generous, and then any that are dually efficient are also more generous towards low and high buyers.

Lemma B.2 (Stacking). *The menus offered in an equilibrium where firms' costs satisfy Assumption 4.2 are partitioned into separate incentive compatibility regions such that*

1. *Low buyers' incentive constraint does not bind in at any menu.*
2. *If some dually efficient menus are offered, there exists a u_i^{de} such that all utilities $u_i < u_i^{de}$ are offered in menus at which high buyers' incentive constraint binds, whereas all utilities $u_i \geq u_i^{de}$ are offered in dually efficient menus.*

Consider the logic that drives this, from the least generous bid to the most generous. The least generous menu is offered by a seller who only expects to sell if it is in a lone match, so it offers exactly the menu it'd choose if it was a monopolist with the same assessment of the buyer's probable valuation, and a monopolist would never offer a menu at which low buyers' incentive constraint binds - featuring $u_h - u_l > q_h^* \Delta \theta$ - since they could strictly increase profits from high sales by offering fewer rents to high buyers. Competition seller types who make more generous offers drives the efficiency of these alongside their generosity (through $u_h - u_l$ growth). When there is sufficient upward pressure on generosity/efficiency, such that the utility gap reaches $u_h - u_l = q_l^* \Delta \theta$, menus become dual efficient. Without an efficiency benefit to high rent concession ($\uparrow u_h$) or efficiency loss to low rent extraction ($\downarrow u_l$) among dually efficient offers, the growth of $u_h - u_l$ slows so that low buyers' incentive constraint never binds.

The results we have covered so far do not rely on the differentiability in any way, but it is a convenient feature to convey intuition and maintain tractability. It is even better to be able to work with continuously differentiable conditional distributions $F_i(u_i|\theta_i)$ over the utilities offered by the average seller to each kind of buyer. Fortunately, equilibria also have these properties.

Lemma B.3 (Equilibrium Distributions). *Equilibrium distributions $F_i(u_i|\theta_i)$ for $i \in \{l, h\}$.*

1. *Do not have atoms in their supports Υ_i .*
2. *Have a convex, connected low support $\Upsilon_l = [\underline{u}_l, \bar{u}_l]$. The high support Υ_h is the union of at most two convex sets disjoint sets, composed of the high utilities offered in menus where high buyers' incentive constraint binds and is slack, respectively. Furthermore, suprema over utilities always satisfy $\bar{u}_i \leq S_i^*$.*
3. *Are continuously differentiable on the interior of their supports with one-sided derivatives at the boundaries.*

Atoms make it possible for sellers to obtain discrete increases in sales in exchange for infinitesimal discounts in profits per sale, so it is clear that these cannot exist. Given that low buyers' incentive constraint does not bind in equilibrium, gaps in the low support would allow the sellers offering a menu with implied utility at the top of the gap to increase their profitability in low matches by

decreasing the rents that the menu offers to low buyers, which would achieve identical low sales but strictly higher profits in each one. The logic for the convexity claim among high offers depends on whether the point u_h that we are considering is such that menus that offer it are dually efficient or constrained. In the former case, high buyers' incentive constraint is slack, so a gap below any such utility would allow a seller that offers it to become more profitable by lowering these high rents, which would preserve high sales, increase high profits per sale, and not affect the efficiency/profitability of its low buyer sales. In the latter case, a constrained menu is ordered when compared with any other, so a gap on the high side is either accompanied by one on the low side (which we have ruled out), an atom at its low utility offer (which we have ruled out), or a situation there are two constrained menus (with the same low utility term, but one has the high utility term at the top of the gap and the other the one at the bottom) and low sales vary locally in low generosity in such a way that it is not preferable to alter the low offer whether the menu is more or less profitable in each low sale, which is not possible. Lastly, the differentiability claims follow because sellers' indifference must be maintained by the probability of winning in combination with profits per sale, and additively separable utilities allowed us to rewrite profits per sale as $S_i(u_l, u_h) - u_i$, which is smooth in marginal changes to either utility, so differentiability of equilibrium distributions becomes necessary to rule out infinitesimal deviations.

B.2 Equilibrium System of Equations Derivation

We now briefly derive the system of equations that allow us to obtain the candidate equilibrium's analytical closed form. Recall that the problem of a type $p_l^{e,j}$ seller is to offer a menu (u_l, u_h) that maximizes her expected profits,

$$\Pi^{e,j}(u_l, u_h) = \sum_{i=l,h} p_i^{e,j} \Psi_i(u_i) (S_i(u_l, u_h) - u_i)$$

subject to the constraint $u_h \geq u_l \geq 0$. The first-order conditions of this problem highlight the interdependence between the optimal amount of utility extended to each type of buyer, as well as the role of posteriors in determining the relative importance of various trade-offs. Based on the fact that only high buyers' incentive constraint can bind in the candidate equilibrium, the seller's optimality conditions are

$$\frac{\partial}{\partial u_l} : \underbrace{p_l^{e,j}(1-\rho)f_l(u_l|\theta_l)(S_l(u_l, u_h) - u_l)}_{\text{sales gains}} - \underbrace{p_l^{e,j}\Psi_l(u_l)}_{\text{rent losses}} + \underbrace{p_l^{e,j}\Psi_l(u_l)\frac{\partial S_l}{\partial u_l}(u_l, u_h)}_{\text{efficiency losses}} = 0 \quad (\text{B.1})$$

$$\frac{\partial}{\partial u_h} : \underbrace{p_h^{e,j}(1-\rho)f_h(u_h|\theta_h)(S_h^* - u_h)}_{\text{sales gains}} - \underbrace{p_h^{e,j}\Psi_h(u_h)}_{\text{rent losses}} + \underbrace{p_l^{e,j}\Psi_l(u_l)\frac{\partial S_l}{\partial u_h}(u_l, u_h)}_{\text{efficiency gains}} = 0 \quad (\text{B.2})$$

Similar terms appear in both equations. The first two capture a typical trade-off between expected sales versus rents per-sale. By increasing indirect utility u_i , a seller makes her offer more attractive to θ_i valuation buyers, thus increasing the probability of selling to them by the mass of equilibrium menus that it would be preferred over in contested matches, mainly $p_i^{e,j}(1-\rho)f_i(u_i|\theta_i)$, whereas the cost of surrendering said rents is directly proportional to the likelihood of trading $p_i^{e,j}\Psi_i(u_i)$ with buyers of this valuation. The third term determines the efficiency effect of an increase in generosity towards θ_i valuation buyers (u_i), and it stems from the point that univariate changes in generosity u_i alter the difference in offered utilities $u_h - u_l$, which drives efficiency. When high buyers' incentive

constraint binds at a menu, generosity towards low buyers (u_l increases) requires further rationing (q_l decrease), thereby reducing the gains of trade with them $S_l(u_l, u_h)$ and, by extension, the profitability of their purchases; the opposite holding for generosity towards high buyers. In other words, generosity in the low (high) offer has an efficiency cost (benefit), when high buyers' incentive constraint is locally binding. Whereas if low buyer's incentive constraint is slack, the efficiency term disappears and the only consideration for the seller is the aforementioned trade-off between the from sales and rents given to buyers of the same valuation.

The implicit objects of immediate interest are the marginal conditional utility distributions $F_i(u_i|\theta_i)$, which shape the nature of competition. Marginals have densities that measure the mass at points on the supports of some seller types' mixed strategy. These supports are atomless, convex, and monotone in the seller's type (overlapping only among any dually efficient bids). Locally, each conditional marginal density $f_i(u_i|\theta_i)$, therefore, corresponds to a weighted density of each seller type's conditional marginal density. In particular, if the utility is offer in an inefficient menu, then there is a unique seller type $p_i^{e,j}$ that offers it and the conditional marginal density is given by,

$$f_i(u_i|\theta_i) = \mu(e)P^e(j|\theta_i)f_i^{e,j}(u_i|\theta_i)$$

For utilities that are offered in dually efficient menus, the conditional marginal density still takes a weighted average form but there is a much simpler way to solve for the utility offers, so we will not use that relation.

To obtain the distribution over utilities in inefficient menus, we will further rewrite the first-order conditions of sellers offering these by applying additional equilibrium properties. In particular, we recall that menus are ordered, so the particular ones of seller types $p_i^{e,j}$ who make inefficient offers are as well, and can be written as functions ($u_l(Q), u_h(Q)$) of the menu's generosity quantile Q in seller type $p_i^{e,j}$'s mixed strategy. This allows us to apply the inverse function theorem and link the conditional marginal densities $f_i(u_i|\theta_i)$ to the progression of utilities,

$$f_i(u_i|\theta_i) = \frac{\mu(e)P^e(j|\theta_i)}{\dot{u}_i^{e,j}(Q)} \quad (\text{B.3})$$

so the conditional marginal distributions take the form,

$$F_i(u_i|\theta_i) = \mu(e)P^e(j|\theta_i)u_i^{e,j,(-1)}(u_i) + \sum_{\substack{e',j' \\ \text{s.t. } p^{e',j'} < p^{e,j}}} \mu(e')P^{e'}(j'|\theta_i) \quad (\text{B.4})$$

where $u_i^{e,j,(-1)}(\cdot)$ is understood to be the inverse of the strictly monotone functions $u_i^{e,j}(Q)$. And, the Q^{th} quantile menu from a $p_i^{e,j}$ type obtains average sales per match,

$$\Psi_i^{e,j}(Q) = \rho + (1 - \rho)F_i(u_i^{e,j}(0)|\theta_i) + (1 - \rho)\mu(e)P^e(j|h)Q$$

in matches with θ_i valuation buyers.

Substituting (B.3) and (B.4) into the first-order conditions produces a standard system of ordinary differential equations that pins down the indirect utilities offered in constrained equilibrium menus. Piecewise linear costs make marginal efficiency effects locally constant, which decouples these equations and allows us to obtain analytical solutions: the equation that drives high utility $u_h^{e,j}(Q)$ is independent, under piecewise linear costs, and we can then substitute its solution into the equation

driving the progression of low utility $u_l^{e,j}(Q)$. Specifically, note the marginal efficiency term becomes,

$$\frac{\partial S_l}{\partial u_h}(u_l, u_h) = (\theta_l - \kappa_l) \frac{\partial q_l}{\partial u_h} = \frac{\theta_l - \kappa_l}{\Delta\theta}$$

So, the differential system governing the progression of utilities in inefficient menus offered by a type $p_l^{e,j}$ seller is,

$$\dot{u}_l^{e,j}(Q) \left[-\frac{\theta_l - \kappa_l}{\Delta\theta} - 1 \right] \Psi_l^{e,j}(Q) + (1 - \rho)\mu(e)P^e(j|l)(S_l(u_l^{e,j}(Q), u_h^{e,j}(Q)) - u_l^{e,j}(Q)) = 0 \quad (\text{B.5})$$

$$\dot{u}_h^{e,j}(Q) \left[\frac{p_l^{e,j}}{p_h^{e,j}} \frac{\theta_l - \kappa_l}{\Delta\theta} \Psi_l^{e,j}(Q) - \Psi_h^{e,j}(Q) \right] + (1 - \rho)\mu(e)P^e(j|h)(S_h^* - u_h^{e,j}(Q)) = 0 \quad (\text{B.6})$$

If offers become dually efficient, either among the offer seller type or because this seller type prefers to jump right to making dually efficient offers when it transitions from those of the adjacent seller type below, we will solve for the remaining utility offers with a different equation. As for the exact form of the utilities that solve the system (B.5)-(B.6), we define $\Xi^{e,j} = \frac{p_h^{e,j}P(j|h)}{p_h^{e,j}P(j|h) - p_l^{e,j}P(j|l)\frac{\theta_l - \kappa_l}{\Delta\theta}}$ to tighten the expressions and write the high utility term as,

$$u_h^{e,j}(Q) = S_h^* - C_h^{e,j} \left(p_h^{e,j} \Psi_h^{e,j}(Q) - p_l^{e,j} \Psi_l^{e,j}(Q) \frac{\theta_l - \kappa_l}{\Delta\theta} \right)^{-\Xi^{e,j}} \quad (\text{B.7})$$

with,

$$C_h^{e,j} = \frac{S_h^* - u_h^{e,j}(0)}{\left(p_h^{e,j} \Psi_h^{e,j}(0) - p_l^{e,j} \Psi_l^{e,j}(0) \frac{\theta_l - \kappa_l}{\Delta\theta} \right)^{-\Xi^{e,j}}} \quad (\text{B.8})$$

which simplifies to,

$$u_h^{e,j}(Q) = S_h^* - (S_h^* - u_h^{e,j}(0)) \left(\frac{p_h^{e,j} \Psi_h^{e,j}(0) - p_l^{e,j} \Psi_l^{e,j}(0) \frac{\theta_l - \kappa_l}{\Delta\theta}}{p_h^{e,j} \Psi_h^{e,j}(Q) - p_l^{e,j} \Psi_l^{e,j}(Q) \frac{\theta_l - \kappa_l}{\Delta\theta}} \right)^{\Xi^{e,j}} \quad (\text{B.9})$$

Substituting this explicit form of high utility into the differential equation for low utility, we then obtain its functional form,

$$u_l^{e,j}(Q) = \frac{\Psi_l^{e,j}(0)}{\Psi_l^{e,j}(Q)} \left(C_l^{e,j} + \frac{(1 - \rho)\mu(e)P^e(j|l)}{\Psi_l^{e,j}(0)} \frac{\theta_l - \kappa_l}{\theta_h - \kappa_l} \left(\frac{C_h^{e,j}}{-p_h^{e,j}P^e(j|h)} \right) \right) \quad (\text{B.10})$$

$$* \left(\left(p_h^{e,j} \Psi_h^{e,j}(Q) - p_l^{e,j} \Psi_l^{e,j}(Q) \frac{\theta_l - \kappa_l}{\Delta\theta} \right)^{1 - \Xi^{e,j}} - \left(p_h^{e,j} \Psi_h^{e,j}(0) - p_l^{e,j} \Psi_l^{e,j}(0) \frac{\theta_l - \kappa_l}{\Delta\theta} \right)^{1 - \Xi^{e,j}} \right) \quad (\text{B.11})$$

with $C_l^{e,j} = u_l^{e,j}(0)$. This heavy expression is not very informative, but we can derive an implicit form of the low utility term - as a function of both its generosity quantile Q and the high utility it is paired with $u_h^{e,j}(\cdot)$ - that is quite helpful. To do this, we simplify (B.6) and rewrite it as,

$$\begin{aligned} \Psi_l^{e,j}(Q) \dot{u}_l^{e,j}(Q) + \psi_l^{e,j}(Q) u_l^{e,j}(Q) &= \psi_l^{e,j}(Q) \frac{\theta_l - \kappa_l}{\theta_h - \kappa_l} u_h^{e,j}(Q) \\ \implies \frac{d}{dQ} \left(\Psi_l^{e,j}(Q) u_l^{e,j}(Q) \right) &= \psi_l^{e,j}(Q) \frac{\theta_l - \kappa_l}{\theta_h - \kappa_l} u_h^{e,j}(Q) \end{aligned}$$

so that,

$$u_l^{e,j}(Q) = u_l^{e,j}(0) \frac{\Psi_l^{e,j}(0)}{\Psi_l^{e,j}(Q)} + \frac{\theta_l - \kappa_l}{\theta_h - \kappa_l} \int_0^Q \frac{\psi_l^{e,j}(x)}{\Psi_l^{e,j}(Q)} u_h^{e,j}(x) dx \quad (\text{B.12})$$

This expression highlights the fact that the low offer is a conditional expectation of the high offers made in less generous menus. This characterization provides an immediate proof for the point that the difference in utilities, $u_h^{e,j}(Q) - u_l^{e,j}(Q)$, and hence efficiency, increases with respect to generosity, and allows us to more easily think about the response of buyer welfare to parameter perturbations, by focusing on the response of high buyer surplus and then averaging these to get that of low buyers.

The initial condition of adjacent types $p_l^{e',j} < p_l^{e',j'}$ depends on whether (a) high buyers' incentive constraint binds at most generous menu $(u_l^{e,j}(1), u_h^{e,j}(1))$ of the lower type $p_l^{e,j}$ and (b) when that happens, whether the next type of seller would prefer for their least generous bid to also be inefficient or dually efficient. This decision is determined by the efficiency gain from additional high rents on the profitability of the least generous low offer,

$$\begin{aligned} & -p_h^{e',j'} \Psi_h^{e',j'}(0) + p_l^{e',j'} \Psi_l^{e',j'}(0)(\theta_l - \kappa_l) \frac{\partial q_l}{\partial u_h} \\ & = -p_h^{e',j'} \Psi_h^{e',j'}(0) + p_l^{e',j'} \Psi_l^{e',j'}(0) \frac{\theta_l - \kappa_l}{\Delta\theta} \end{aligned}$$

When this is > 0 , sellers of the higher type $p_l^{e',j'}$ prefer for their lowest bid to be dually efficient, and there is a discontinuous jump in generosity between $u_h^{e,j}(1) < u_h^{e',j'}(0)$. When the condition is < 0 , instead, sellers of the higher type prefer for their least generous offers to also be inefficient, and there is no discontinuity between $u_h^{e,j}(1) < u_h^{e',j'}(0)$. In either case, the progression of the low utility terms is continuous and the higher seller type's least generous low utility offer is exactly the most generous one of the lower seller type, $u_l^{e,j}(1) = u_l^{e',j'}(0)$.

We will prove that the difference in utilities $u_h^{e,j}(Q) - u_l^{e,j}(Q)$ increases in the seller's quantile and hence the efficiency of the offer, so if these reach the point $u_h^{e,j}(Q) - u_l^{e,j}(Q) = q_l^* \Delta\theta$ such that the menu becomes dually efficient, then we solve for the dually efficient menus with a simpler equation. In particular, given the utility pair $(u_l^{de}(0), u_h^{de}(0))$ at which menus become dually efficient and the mass of sellers that makes constrained offers in θ_i matches $F_i(u_i^{de})$, then the equations that low and high offers must satisfy for dually efficient bids to have the necessary property of being equally profitable in low (high) matches is,

$$(\rho + (1-\rho)F_i(u_i^{de}(0)|\theta_i))(S_i^* - u_i^{de}(0)) = (\rho + (1-\rho)(1 - F_i(u_i^{de}(0)|\theta_i))Q)(S_i^* - \underline{u}_i^{de}(Q)) \quad \text{for } i \in \{l, h\} \quad (\text{B.13})$$

where Q is the quantile among dually efficient menus of the utility $u^{de}(Q)$, and $1 - F_i(u_i^{de}(0)|\theta_i)$ is the probability that a seller makes a dually efficient offer in a θ_i match. By construction, these indirect utility functions have all the properties stipulated earlier: strictly increasing in generosity, monotone in seller type, identically ranked by their low and high utility offerings, and jointly forming a bottom pair of intervals comprised of utilities offered in menus where high buyers' incentive constraint binds, potentially followed above by another pair of intervals comprised of utilities offered in dually efficient menus.

Appendix C Endogenous Quality: Equilibrium Property Proofs

Proof of Theorem 4.1. Per sale profits from θ_i valuation buyer intended contracts take the form:

$$\pi(q_i, x_i) = x_i - \phi(q_i) = (\theta_i q_i - \phi(q_i)) + (x_i - \theta_i q_i) = S_i(q_i) - u_i \quad (\text{C.1})$$

and $S_i(q_i)$ is concave, reaching a maximum at q_i^* . Due to concavity, any incentive-compatible menu featuring a low quality $q_l > q_l^*$ is strictly dominated by one with a revised low contract of $(q_l^*, u_l -$

$\theta_l(q_l - q_l^*)$), while any incentive-compatible menu featuring a high quality $q_h < q_h^*$ is strictly dominated by one with a revised high contract of $(q_h^*, u_h + \theta_h(q_h^* - q_h))$. In other words, the optimal incentive-compatible menus feature $q_l \leq q_l^*$ and $q_h \geq q_h^*$.

Furthermore, IC_i must bind¹⁹ if the θ_{-i} intended contract is not efficient; else, the revised contract $(\tilde{q}_{-i}, \tilde{x}_{-i})$ featuring²⁰ $\tilde{q}_{-i} = \frac{u_h - u_l}{\Delta\theta}$ and $\tilde{x}_{-i} = \theta_{-i}\tilde{q}_{-i} - u_{-i}$ could be paired with the former θ_i contract for a strictly dominant menu - same expected θ_{-i} sales (by preserving the utility that a θ_{-i} valuation buyer obtains), strictly higher profits-per-sale in θ_{-i} matches (by preserving u_{-i} and increasing the gains from trade (see (C.1))), and maintaining the incentive compatibility of θ_i buyers ($\theta_i\tilde{q}_{-i} - \tilde{x}_{-i} = u_i$). Inversely, since only one constraint IC can bind in a given menu, profit maximality implies that any menu featuring an inefficient q_{-i} offer, also features efficient $q_i = q_i^*$ quality provision.

Incentive compatibility bounds the qualities offered in each contract by the suggested ratio: $q_l \leq \frac{u_h - u_l}{\Delta\theta}$ (to satisfy IC_h) and $q_h \geq \frac{u_h - u_l}{\Delta\theta}$ (to satisfy IC_l). We have shown that IC_i binds when $q_{-i} \neq q_i^*$ though, so equality of the respective bound yields the form of the inefficient quality $q_{-i} = \frac{u_h - u_l}{\Delta\theta}$. Lastly, any menu featuring efficient quality provision in both contracts must feature utility offerings satisfying $q_l^*\Delta\theta \leq u_h - u_l \leq q_h^*\Delta\theta$.

We conclude that: (1) menus featuring inefficient low valuation buyer provision are IC_h binding, feature $q_l = \frac{u_h - u_l}{\Delta\theta} < q_l^*$, and are paired with efficient high valuation buyer contracts, (2) menus featuring inefficient high valuation buyer provision are IC_l binding, feature $\frac{u_h - u_l}{\Delta\theta} = q_h < q_h^*$, and are paired with efficient low valuation buyer contracts, and (3) doubly efficient menus correspond to those for which $q_l^*\Delta\theta \leq u_h - u_l \leq q_h^*\Delta\theta$ almost all of which have locally slack IC constraints with the exception of boundary ones satisfying $q_i^*\Delta\theta = \frac{u_h - u_l}{\Delta\theta}$ at which IC_{-i} binds. \square

Proof of Theorem 4.2. Recall that seller types with degenerate beliefs have measure zero, so the following arguments apply to almost every bid offered in a match, in particular those that would be offered by types with nondegenerate conditional beliefs.

Consider the essential infimum and supremum, \underline{u}_i and \bar{u}_i , respectively, on the equilibrium indirect utility offered to θ_i valuation buyers. By Theorem 4.1, equilibrium menus are separating and these must correspond to the bounds on utilities extended in θ_i intended contracts, with profits $S_i(u_l, u_h) - u_i$. If $\bar{u}_i > S_i^*$, then a discrete mass of menus would feature contracts with a non-zero probability of being accepted and entail negative profits per sale. As such, $\bar{u}_i \leq S_i^*$ and we will argue that the lower bound \underline{u}_i is exactly S_i^* .

Suppose that $\underline{u}_i < S_i^*$. By Theorem 4.1, optimal contracts with $u_i < S_i^*$ entail strictly positive²¹ profits per sale if $q_i > 0$ and zero profits per sale if $q_i = 0$. Further, there are only two possibilities in a neighborhood of \underline{u}_i : either \underline{u}_i is an atom, or half-open sets $[\underline{u}_i, \underline{u}_i + \delta)$ are assigned arbitrarily small mass, as $\delta \searrow 0$, by the equilibrium distribution of indirect utility offerings by competitors in θ_i matches $F_i(u_i|\theta_i)$.

Recall that every contract makes nonnegative profits per sale, since not trading (offering zero quality) is always an option. As such, if there were an atom at \underline{u}_i , the menu implied by pairing a slightly more generous θ_{-j} offering $u_{-j} + 2\varepsilon$ and $\underline{u}_i + \varepsilon$ (for $\varepsilon > 0$ small) would be strictly dominant, entailing (at worst) an arbitrarily small decrease in θ_{-j} per sale profits, a discrete increase in θ_i expected sales, and thus a discrete increase in θ_i profits per match. Whereas if there were no atom at the lower bound, contracts offering u_i arbitrarily close have θ_i match profits arbitrarily close to zero

¹⁹It is elementary to check that only one IC constraint can bind in a menu.

²⁰Where $\tilde{q}_{-i} < q_{-i}$ if $-i = l$ and $>$ if $-i = h$.

²¹If $q_i = q_i^*$, the claim follows. If $q_i \neq q_i^*$ and $u_i = S_i(q_i)$, then lowering u_i by a small $\varepsilon < u_i - \underline{u}_i$ would allow the seller to still win matches with some probability u_i and make strictly positive profits in these, as the revision would both increase the efficiency and lower the generosity of these sales.

(they almost surely compete against a more generous seller); so a discrete mass of these is strictly dominated by a menu of the form given in the atom case.

We conclude that almost every seller selects a pair of contracts that extend implied utilities (S_l^*, S_h^*) . By Theorem 4.1, the unique optimal menu that satisfies these conditions is $((q_l^*, \phi(q_l^*)), (q_h^*, \phi(q_h^*)))$. \square

No Atoms. Toward a contradiction, suppose that F_h had an atom at u_h and let (u_l, u_h) be an equilibrium menu featuring this high bid generosity.

We begin by showing that $S_h(u_l, u_h) - u_h > 0$ in any equilibrium offer (u_l, u_h) . Suppose not. Then we must have $S_l(u_l, u_h) - u_l \leq 0$, otherwise offering a pooling menu of only the low valuation buyer's contract would strictly increase the seller's expected profits - strictly positive per-sale profits from θ_h sales and strictly positive probability of being accepted ($\rho > 0$), even by low buyers. But then expected profits $\Pi \leq 0$, which contradicts seller optimization: the seller can always offer the menu $((0, 0), (q_h^*, \theta_h q_h^*))$ and obtain strictly positive expected profits. With this fact, we can rule out an atom at any u_h in $\text{supp}(\Psi_h)$.

In particular, note that at any u_h with discrete mass and for any type of seller p_l ,

$$\begin{aligned} & \lim_{\varepsilon \searrow 0} \Pi(u_l + \varepsilon, u_h + \varepsilon) - \Pi(u_l, u_h) \\ &= \lim_{\varepsilon \searrow 0} \left\{ \sum_{k=l,h} p_k \Psi_k(u_k + \varepsilon) (S_k(u_l + \varepsilon, u_h + \varepsilon) - u_k - \varepsilon) - \sum_{k=l,h} p_k \Psi_k(u_k) (S_k(u_l, u_h) - u_k) \right\} \\ &= \lim_{\varepsilon \searrow 0} ((1 - \rho)(F_h(u_h + \varepsilon) - F_h(u_h)))(S_h(u_l, u_h) - u_h) \\ &> 0 \end{aligned} \tag{C.2}$$

and so (u_l, u_h) would be strictly dominated by $(u_l + \varepsilon, u_h + \varepsilon)$, for some $\varepsilon > 0$.

Furthermore, $S_l(u_l, u_h) - u_l \geq 0$ in equilibrium; otherwise, this menu would be strictly dominated by one with the paired offers $(q_l, x_l) = (0, 0)$ and $(q_h, x_h) = (q_h^*, \theta_h q_h^* - u_h)$, which maintain expected high match profits and strictly those of low ones (these buyers either select a no-loss contract, $(0, 0)$, or one with strictly positive profits per sale, (q_h, x_h)).

We close by ruling out atoms among low bids. Suppose that $\text{supp}(\Psi_l)$ had an atom at u_l and let (u_l, u_h) be an equilibrium menu featuring this low bid generosity. Inequalities (C.2) rule out $S_l(u_l, u_h) - u_l > 0$, so the only possible menu with such a low offering must be one that makes zero profits in low matches $S_l(u_l, u_h) - u_l = 0$. Suppose that $u_l > 0$. If IC_l is slack, then menu with slightly less low generosity $(u_l - \varepsilon, u_h)$ is strictly dominant (low efficiency nondecreasing hence positive profits per low sale, some low sales since $\rho > 0$, high trade profitability not affected), whereas if IC_l binds, low buyers obtain utility $u_l > 0$ from the high contract, which has strictly positive profits per sale, so the seller could just pool all buyers on this contract and makes strictly positive profits in low matches as well. Lastly, if $u_l = 0$ and $S_l(u_l, u_h) - u_l = 0$, it follows (by Theorem 4.1) that $u_h = 0$, so a strictly positive mass of such (u_l, u_h) menus would give rise to an atom in $\text{supp}(\Psi_h)$. \square

Weak-Orderedness. Consider two equilibrium menus (u_l, u_h) and $(\tilde{u}_l, \tilde{u}_h)$, offered by sellers of (possibly equal) respective types p_l and \tilde{p}_l , which violate weak-orderedness; without loss, suppose that this takes place via $\tilde{u}_h > u_h$ and $u_l > \tilde{u}_l$. We proceed case-by-case, depending on the incentive compatibility constraint that binds at (u_l, u_h) .

Suppose that IC_h binds at (u_l, u_h) . Then,

$$\begin{aligned}\Psi_l(u_l) &> \Psi_l(\tilde{u}_l) \\ S_l(u_l, \tilde{u}_h) - S_l(u_l, u_h) &\geq S_l(\tilde{u}_l, \tilde{u}_h) - S_l(\tilde{u}_l, u_h) \geq 0 \\ S_l(u_l, \tilde{u}_h) - S_l(u_l, u_h) &> 0\end{aligned}$$

where the second set of inequalities follows by the convexity of costs, Theorem 4.1,

$$\begin{aligned}S_l(u_l, u_h) &= \theta_l q_l(u_l, u_h) - \phi(q_l(u_l, u_h)) \\ q_l(u_l, u_h) &= \begin{cases} \frac{u_h - u_l}{\Delta\theta} & \text{if } u_h - u_l < q_l^* \Delta\theta \\ q_l^* & \text{otherwise} \end{cases}\end{aligned}$$

while the third is due to the fact that IC_h binds at (u_l, u_h) (by assumption), so the efficiency of low trade achieved by (u_l, \tilde{u}_h) must be strictly greater. Thus,

$$\Psi_l(u_l) (S_l(u_l, \tilde{u}_h) - S_l(u_l, u_h)) - \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, \tilde{u}_h) - S_l(\tilde{u}_l, u_h)) > 0$$

and,

$$\Psi_l(u_l) (S_l(u_l, \tilde{u}_h) - u_l) - \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_l) > \Psi_l(u_l) (S_l(u_l, u_h) - u_l) - \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, u_h) - \tilde{u}_l) \quad (C.3)$$

When facing high valuation buyers, Theorem 4.1 indicates that,

$$\begin{aligned}S_h(u_l, u_h) &= \theta_h q_h(u_l, u_h) - \phi(q_h(u_l, u_h)) \\ q_h(u_l, u_h) &= \begin{cases} \frac{u_h - u_l}{\Delta\theta} & \text{if } u_h - u_l > q_h^* \Delta\theta \\ q_h^* & \text{otherwise} \end{cases}\end{aligned}$$

so high contracts are decreasing in efficiency with respect to the difference in utility that a menu offers to high vs low valuation buyers. But the weak-orderedness violation implies $\tilde{u}_h - u_l < \tilde{u}_h - \tilde{u}_l$ and $u_h - u_l < u_h - \tilde{u}_l$, so that cost convexity also implies,

$$S_h(u_l, \tilde{u}_h) - S_h(\tilde{u}_l, \tilde{u}_h) \geq S_h(u_l, u_h) - S_h(\tilde{u}_l, u_h) \geq 0$$

which is equivalent to,

$$\Psi_h(\tilde{u}_h) ((S_h(u_l, \tilde{u}_h) - \tilde{u}_h) - (S_h(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_h)) \geq \Psi_h(u_h) ((S_h(u_l, u_h) - u_h) - (S_h(\tilde{u}_l, u_h) - u_h)) \geq 0 \quad (C.4)$$

Jointly (C.3) and (C.4) yield,

$$\begin{aligned}& p_l [\Psi_l(u_l) (S_l(u_l, \tilde{u}_h) - u_l) - \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_l)] \\ & + p_h [\Psi_h(\tilde{u}_h) (S_h(u_l, \tilde{u}_h) - \tilde{u}_h) - \Psi_h(\tilde{u}_h) (S_h(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_h)] \\ & > p_l [\Psi_l(u_l) (S_l(u_l, u_h) - u_l) - \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, u_h) - \tilde{u}_l)] \\ & + p_h [\Psi_h(u_h) (S_h(u_l, u_h) - u_h) - \Psi_h(u_h) (S_h(\tilde{u}_l, u_h) - u_h)] \\ & \geq 0\end{aligned}$$

where the last inequality comes from the optimality of (u_l, u_h) for a p_l type seller. Necessarily then, at least one of the terms in brackets at the topmost expression must be > 0 and we know, by (C.4),

that the second term is ≥ 0 . If the first was > 0 or if it were $= 0$ (and so the second term in brackets was > 0), then any seller - including a \tilde{p}_l type - would strictly prefer (u_l, \tilde{u}_h) over $(\tilde{u}_l, \tilde{u}_h)$. The only alternative then is that this first bracket term is < 0 and so that the second bracket term is > 0 . Given this fact, we can take advantage of (C.3), to note that the first term in brackets of the bottom expression is also < 0 and the second > 0 . The latter strict inequality can only hold when IC_l binds at both (\tilde{u}_l, u_h) and, since $\tilde{u}_h > u_h$, also at $(\tilde{u}_l, \tilde{u}_h)$.

However, for (u_l, u_h) to be preferred by a type p_l seller over (u_l, \tilde{u}_h) , which is strictly more profitable in low matches (since (u_l, u_h) is IC_h binding), the chosen menu must be strictly more profitable in the high. But then,

$$\begin{aligned} & \Psi_h(u_h) (S_h(\tilde{u}_l, u_h) - u_h) - \Psi_h(\tilde{u}_h) (S_h(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_h) \\ & \geq \Psi_h(u_h) (S_h(u_l, u_h) - u_h) - \Psi_h(\tilde{u}_h) (S_h(u_l, \tilde{u}_h) - \tilde{u}_h) \\ & > 0 \end{aligned} \tag{C.5}$$

and (\tilde{u}_l, u_h) is strictly better than $(\tilde{u}_l, \tilde{u}_h)$ in high matches and identical in the low (same sales, same profits per-sale) - making it strictly preferable for a \tilde{p}_l type.

These arguments are sufficient to establish that²² for a given optimal menu (u_l, u_h) at which IC_h binds, then any other equilibrium menu $(\tilde{u}_l, \tilde{u}_h)$ with $\tilde{u}_i > u_i$ - offered by any type of seller - must have $\tilde{u}_{-i} \geq u_{-i}$.

The only comparisons left to consider are those between two IC_l binding menus and those between one at which IC_l binds one with another that is dually efficient; any comparison where IC_h binds in a menu is covered by the previous reasoning. But, if IC_l is to bind in either of the potentially non-weakly-ordered menus, giving rise to $\tilde{u}_h > u_h$ and $u_l > \tilde{u}_l$, then it must bind at the menu with the largest utility difference, $(\tilde{u}_l, \tilde{u}_h)$. Suppose that this is so and consider the alternative bid (\tilde{u}_l, u_h) .

By the fact that (u_l, u_h) is either dually efficient or IC_l binding, the same must be true for (\tilde{u}_l, u_h) and (u_l, \tilde{u}_h) , which have strictly larger differences in offered utilities, so $S_l^* = S_l(u_l, u_h) = S_l(\tilde{u}_l, u_h) = S_l(u_l, \tilde{u}_h) = S_l(\tilde{u}_l, \tilde{u}_h)$ and,

$$\begin{aligned} & \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, u_h) - \tilde{u}_l) - \Psi_l(\tilde{u}_l) (S_l(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_l) \\ & = \Psi_l(u_l) (S_l(u_l, u_h) - u_l) - \Psi_l(u_l) (S_l(u_l, \tilde{u}_h) - u_l) \\ & = 0 \end{aligned} \tag{C.6}$$

Given that IC_l binds at $(\tilde{u}_l, \tilde{u}_h)$ however, $S_h(u_l, \tilde{u}_h) - S_h(\tilde{u}_l, \tilde{u}_h) > S_h(u_l, u_h) - S_h(\tilde{u}_l, u_h)$ and,

$$\begin{aligned} & \Psi_h(u_h) (S_h(\tilde{u}_l, u_h) - u_h) - \Psi_h(\tilde{u}_h) (S_h(\tilde{u}_l, \tilde{u}_h) - \tilde{u}_h) \\ & > \Psi_h(u_h) (S_h(u_l, u_h) - u_h) - \Psi_h(\tilde{u}_h) (S_h(u_l, \tilde{u}_h) - \tilde{u}_h) \\ & \geq 0 \end{aligned} \tag{C.7}$$

where the weak inequality is due to the choice of (u_l, u_h) over (u_l, \tilde{u}_h) by the p_l type seller, while the strict inequality follows from the IC_l constraint binding in all four of menus under consideration. Jointly (C.6) and (C.7) then imply that $(\tilde{u}_l, \tilde{u}_h)$ is strictly dominated by (\tilde{u}_l, u_h) for any seller type however, since the latter is equally profitable in low matches and strictly better in the high. \square

Ordered, Support Convexity, and Profit Ranking. We will now strengthen the claim that equilibrium

²²Note that the case of IC_h binding at (u_l, u_h) and $u_h > \tilde{u}_h$ but $\tilde{u}_l > u_l$ is subsumed in one we've established, because IC_h must also bind at $(\tilde{u}_l, \tilde{u}_h)$ in these other inequalities.

menus are weakly ordered to one of (strict) orderedness. The proof is broken up into 8 steps which will show the monotonicity of profits per low and high match with respect to generosity as well as convex supports for the probability of winning distributions Ψ_j .

Step 1 (Conditional Expected Sales): *Every type of seller expects the same sales, $\Psi_i(u_i)$ from an indirect utility offer of u_i .*

The buyer's type determines the distribution of signals observed by sellers of each expertise,

$$P^e(j|\theta_i)$$

and so the distribution of seller types,

$$P(p_l = p|\theta_i) = \sum_{e, j} \mathbb{1}(p = P^e(\theta_l|j)) P^e(j|\theta_i) \mu(e)$$

A seller's mixing distribution, $P((u_l, u_h)|p_l)$, is determined by her type. As such, every seller's expected distribution of competitor bids,

$$F(\tilde{u}_i \leq u_i|\theta_i) = \sum_{p_l} P(\tilde{u}_i \leq u_i|p) P(p_l|\theta_i)$$

is identical when the buyer's type is θ_i . \square

Step 2 (No IC_l): *Under Assumption 4.2, IC_l binding menus are not offered in equilibrium.*

The profits per sale high sale of any menu (u_l, u_h) are given by,

$$S_h(u_l, u_h) - u_h \leq S_h^* - u_h$$

so $u_h \leq S_h^*$, as losses are strictly increasing in u_h for larger values. The gap between utilities of any menu is therefore bounded by,

$$\begin{aligned} u_h - u_l &\leq u_h \leq S_h^* = (\theta_h - \kappa_m)(q_h^* - q_l^*) + (\theta_h - \kappa_l)q_l^* \\ &\leq (\theta_h - \theta_l)(q_h^* - q_l^*) + (\theta_h - \theta_l)q_l^* \\ &= \Delta\theta q_h^* \end{aligned}$$

Step 3 (Dually Efficient Characterization): *If $\Delta\theta q_l^* < \Delta u < \Delta\theta q_h^*$, then sufficiently small changes u_h or u_l leave expected profits per low and high match unchanged. Further, the existence of a menu $(\bar{u}_l^{de}, \bar{u}_h^{de})$ with $\Delta\theta q_l^* < \bar{u}_h^{de} - \bar{u}_l^{de} \leq \Delta\theta q_h^*$ implies the existence of another menu $(\underline{u}_l^{de}, \underline{u}_h^{de})$ with $\underline{u}_h^{de} - \underline{u}_l^{de} = \Delta\theta q_l^*$ and convex regions $[u_i^{de}, \bar{u}_i^{de}]$ for $i \in \{l, h\}$ made up of offerings from dually efficient menus that satisfy the strict inequalities. Lastly, if $\Delta\bar{u}^{de} = \Delta\theta q_h^*$, then there exists a $\delta > 0$, such that a bid of $(u_l - \delta, u_h - \delta)$ has identical profits in low and high matches.*

Since lowering either offering does not affect the efficiency of bids when the inequalities are strict, a gap below either u_i would allow a strict increase in profitability from undercutting the original utility by some $\epsilon > 0$. Given that there aren't gaps under either coordinate then, there can't be IC_h binding menus immediately below - by weak-orderedness and these menus' strictly smaller Δu respectively - so, all coordinates $u'_i \in [u_i - \delta, u_i]$ for $\delta > 0$

small and $j \in \{l, h\}$ belong to dually efficient menus, and consequently, $\Psi_i(u'_i)(S_i^* - u'_i) = \Psi_i(u_i)(S_i^* - u_i)$ (otherwise, some seller would be able to deviate to a bid with higher expected profits per-match without affecting the efficiency of the paired contract). This establishes the first sentence's claim.

The preceding logic implies that dually efficient menus satisfying the pair of strict inequalities have other dually efficient menus immediately below them in generosity and that the profitability of offering these menus, conditionally on matching with either type of buyer, is identical. Consider the infimum u_i coordinates of the contiguous intervals made up by these dually efficient menus and let us refer to it as \underline{u}_i^{de} for $i \in \{l, h\}$. Note that $\Delta \underline{u}^{de} = \Delta \theta q_l^*$, since otherwise necessarily there'd be a gap below one of the \underline{u}_i^{de} coordinates (no sufficiently close dually efficient menus with coordinates below, whereas coordinates from IC_h binding menus would create gaps), which would allow an analogous deviation²³ as in the previous paragraph for those seller types offering dually-efficient menus with u_i 's close to \underline{u}_i^{de} . This establishes the second sentence.

For the case that $\Delta \bar{u}^{de} = \Delta \theta q_h^*$, it is sufficient to observe that all entries immediately below the menu $(\bar{u}_l^{de}, \bar{u}_h^{de})$ and sufficiently close must also be dually efficient to avoid creating a gap. By continuing down distribution in either coordinate one must eventually reach a dually efficient menu with $\Delta u < \Delta \theta q_h^*$; else, there'd either be a gap (allowing the familiar deviation) if one encountered an IC_h menu before reaching a dually efficient one that satisfied the pair of strict inequalities, or it'd hold that $\underline{u}_l^{de} = 0$ with a paired high utility of $\tilde{u}_h = \underline{u}_l^{de} + \Delta \theta q_h^*$. In this last case, the lack of atoms, weak-orderedness (ruling out a non-zero mass of IC_h binding ones having entries below \tilde{u}_h), and lack of dually efficient bids with $\Delta u < \Delta \theta q_h^*$ would imply that there is a gap below \tilde{u}_h and so that this menu's profitability could be strictly improved by lowering its high offering (more profits per high sale without decreasing the efficiency of low sales).

To prove the statement in the last sentence, consider the supremum over dually efficient bids u_i with $u_i < \bar{u}_j^{de}$ that belong to a menu with $\Delta u < \Delta \theta q_h^*$, and refer to it as u'_j . If $u'_j = \bar{u}_j^{de}$, we are done by the previous claims. If $u'_j < \bar{u}_j^{de}$, then all menus with bids in $[u'_j, \bar{u}_j^{de}]$ for either $j \in \{l, h\}$ must have $\Delta u = \Delta \theta q_h^*$ and it is sufficient to establish the invariance of profits in this region, as the previous arguments take over for dually efficient menus with bids below. Note that the existence of a high offering u_h that allows greater expected profits per high match than neighboring ones above would allow these sellers above to strictly increase their expected profits per high match (while leaving low match ones unperturbed), by instead choosing u_h as the high pairing; by extension of this local argument, expected high match profits are maintained by menus satisfying $\Delta \theta q_h^*$ in the interval $[u'_h, \bar{u}_h^{de}]$. The invariance of expected low match profits for menus with bids in $[u'_l, \bar{u}_l^{de}]$ follows by a similar logic: the existence of a point u_l in this interval allowing larger expected low match profits would allow sellers above to shift their high and low offering by the same amount so as to obtain the superior expected profits per low match while preserving the same expected profits per high match.

Step 4 (Convex Ψ_l Support): *There are no gaps in $\text{supp}(\Psi_l)$.*

Since IC_l never binds in an equilibrium bid by Step 2, gaps in $\text{supp} \Psi_l$ would imply the existence of a deviation for sellers bidding menus featuring a u_l offer close to the top of

²³Profits conditionally on buyer type are continuous with respect to generosity due to the lack of atoms in the bid distribution, while the continuity of profits per sale follows from the efficiency formulation $S_i(u_l, u_h) - u_i$.

said gap. Such a seller could decrease its low offer to some value in the gap, so as to obtain a discrete increase in profits per low sale, while sacrificing an arbitrarily small number of low sales (sales continuous in generosity) and maintaining expected profits per high match, thereby strictly increasing expected profits. The only non-obvious case is if $\Delta u = \Delta \theta q_h^*$, but then we know from Step 3 that all the u'_h immediately below u_h preserve expected profits per high match, so a revision of u_l and u_h by the same δ downward would do what is described.

Step 5 (Ordereness Violations): *The only possible violation of strict-orderedness between two equilibrium menus (u_l, u_h) and (u'_l, u'_h) , where IC_h binds in at least one of them, is if $u_l = u'_l$, $u'_h > u_h$, and $p'_l > p_l$ for the respective seller types that offer these.*

Given two equilibrium menus (u_l, u_h) and (u'_l, u'_h) , such that IC_h binds in at least one, these must be weakly ordered, so the violation of strict ordering must involve either $u_l = u'_l$ and $u_h > u'_h$ or $u_h = u'_h$ and $u_l > u'_l$. However, the second case cannot be.

To see this point, note that if $u_h = u'_h$ and $u_l > u'_l$, IC_h must bind at the menu with the smaller generosity difference, (u_l, u_h) . As such, any other bid $(\tilde{u}_l, \tilde{u}_h)$ with $\tilde{u}_l \in [u'_l, u_l]$ is weakly ordered with respect to this menu, so $\tilde{u}_h \leq u_h$. And if IC_h bound at $(\tilde{u}_l, \tilde{u}_h)$, then it too would be weakly ordered with respect to (u'_l, u'_h) , so that $u'_h \leq \tilde{u}_h$. Jointly, these statements imply that $u_h = \tilde{u}_h = u'_h$ for any IC_h binding bid with a low offering in interval $[u'_l, u_l]$ and thus that a nonzero mass of such menus gives rise to an atom at u_h (contradicting the lack of atoms). We only have the mass of dually efficient menus left to fill the intervals below u_l . Weak ordering requires these to satisfy $\tilde{u}_h \leq u_h$ and dual efficiency $\Delta \theta q_l^* \leq \Delta \tilde{u}$, so that,

$$\tilde{u}_l \leq \tilde{u}_h - \Delta \theta q_l^* \leq u_h - \Delta \theta q_l^* < u_l$$

Implying the existence of a $\delta > 0$, for which $\Psi_l([u_l - \delta, u_l]) = 0$ (contradicting Step 4).

If $u'_l = u_l$ and $u'_h > u_h$ instead, IC_h must bind at (u_l, u_h) ; therefore, (u'_l, u'_h) is strictly more profitable in low matches (more efficient, same low generosity, same sales). The menu (u_l, u_h) is offered in equilibrium though, so it cannot be strictly dominated and must be superior to (u'_l, u'_h) in high matches. This difference in profits conditionally on a buyer type can only be optimal for two sellers if they differ in seller type, $p'_l \neq p_l$, with the one who places more weight on low matches $p'_l > p_l$ offering the menu (u'_l, u'_h) that is more profitable in them.

Step 6 (Profit Ranking): *Given two menus (u_l, u_h) and (u'_l, u'_h) with $u_l < u'_l$, it must be that the expected profits per high match satisfy $\Psi_h(u'_h)(S_h^* - u'_h) \leq \Psi_h(u_h)(S_h^* - u_h)$ and inversely for low match profits, $\Psi_l(u'_l)(S_l(u'_l, u'_h) - u'_l) \geq \Psi_l(u_l)(S_l(u_l, u_h) - u_l)$, with both inequalities either strict or equal. In particular, if the inequalities are strict, the type p'_l of the seller who bids (u'_l, u'_h) must be strictly greater than that of the seller who bids (u_l, u_h) .*

Consider menus (u_l, u_h) and (u'_l, u'_h) with $u_l < u'_l$ close. If $\Delta \theta q_l^* \leq \Delta u' \leq \Delta \theta q_h^*$, any (u_l, u_h) with u_l sufficiently close must be dually efficient and maintain expected profits per high and low match by Step 3. Whereas if $\Delta u' < q_l^* \Delta \theta$, for u_l close: (a) $\Pi_h(u'_l, u'_h) = \Pi_h(u_l, u'_h)$, (b) to rule out a deviation toward the more efficient bid (u_l, u'_h) by the type bidding (u_l, u_h) , necessarily $\Pi_h(u_l, u'_h) < \Pi_h(u_l, u_h)$, and (c) to rule out a deviation, by

the type offering (u'_l, u'_h) , towards the more high match profitable bid (u_l, u_h) , necessarily $\Pi_l(u'_h, u'_l) > \Pi_l(u_h, u_l)$. The menu with the more generous low offer is therefore more profitable in low matches and must be offered by the seller of type p_l , which places more weight on low matches. This local reasoning around any point $\text{supp}(\Psi_l)$ yields the global claim. □

Continuous Differentiable Distributions. We will show that the functions $\Psi_h(\cdot)$ and $\Psi_l(\cdot)$ are continuously differentiable. Since these functions are given by the composition of F_h and F_l with a continuous mononote function, this will prove the continuous differentiability of the equilibrium offer distributions.

We present the case of Ψ_h (Ψ_l 's is analogous). Let u_h be a utility level offered in the interior of the support of F_h and u_l be its paired low utility offering such that the menu (u_l, u_h) is optimal for some seller of type p_l . We proceed to bound the difference $\Psi_h(u_h + \varepsilon) - \Psi_h(u_h)$ from above and below.

For any $\varepsilon \in \mathbb{R}$, the expected profit to this type of seller from the deviation menu $(u_l, u_h + \varepsilon)$ can be decomposed as:

$$\begin{aligned} & p_l \Psi_l(u_l)(S_l(u_l, u_h + \varepsilon) - u_l) + p_h \Psi_h(u_h + \varepsilon)(S_h(u_l, u_h + \varepsilon) - u_h - \varepsilon) \\ &= p_l \Psi_l(u_l)(S_l(u_l, u_h) - u_l) + p_h \Psi_h(u_h)(S_h(u_l, u_h) - u_h) \\ &+ p_l \Psi_l(u_l)(S_l(u_l, u_h + \varepsilon) - S_l(u_l, u_h)) + p_h \Psi_h(u_h)(S_h(u_l, u_h + \varepsilon) - \varepsilon - S_h(u_l, u_h)) \\ &+ p_h(\Psi_h(u_h + \varepsilon) - \Psi_h(u_h))(S_h(u_l, u_h + \varepsilon) - u_h - \varepsilon) \end{aligned}$$

and by optimality, we must have $\Pi(u_l, u_h; p_l) \geq \Pi(u_l, u_h + \varepsilon; p_l)$, which implies the following inequality,

$$\begin{aligned} & p_h(\Psi_h(u_h + \varepsilon) - \Psi_h(u_h))(S_h(u_l, u_h + \varepsilon) - u_h - \varepsilon) \\ & \leq p_l \Psi_l(u_l)(S_l(u_l, u_h) - S_l(u_l, u_h + \varepsilon)) + p_h \Psi_h(u_h)(S_h(u_l, u_h) - S_h(u_l, u_h + \varepsilon) + \varepsilon) \end{aligned} \tag{C.8}$$

Similarly, for any $\varepsilon \in \mathbb{R}$ such that $u_h + \varepsilon$ is in the interior of the support of F_h , there exists $u_{l,\varepsilon}$ for which $(u_{l,\varepsilon}, u_h + \varepsilon)$ is the optimal bid of some seller with type \tilde{p}_l . And, decomposing the profits from this bid to said seller as above:

$$\begin{aligned} & \tilde{p}_l \Psi_l(u_{l,\varepsilon})(S_l(u_{l,\varepsilon}, u_h + \varepsilon) - u_{l,\varepsilon}) + \tilde{p}_h \Psi_h(u_h + \varepsilon)(S_h(u_{l,\varepsilon}, u_h + \varepsilon) - u_h - \varepsilon) \\ &= \tilde{p}_l \Psi_l(u_{l,\varepsilon})(S_l(u_{l,\varepsilon}, u_h) - u_{l,\varepsilon}) + \tilde{p}_h \Psi_h(u_h)(S_h(u_{l,\varepsilon}, u_h) - u_h) \\ &+ \tilde{p}_l \Psi_l(u_{l,\varepsilon})(S_l(u_{l,\varepsilon}, u_h + \varepsilon) - S_l(u_{l,\varepsilon}, u_h)) + \tilde{p}_h \Psi_h(u_h)(S_h(u_{l,\varepsilon}, u_h + \varepsilon) - \varepsilon - S_h(u_{l,\varepsilon}, u_h)) \\ &+ \tilde{p}_h(\Psi_h(u_h + \varepsilon) - \Psi_h(u_h))(S_h(u_{l,\varepsilon}, u_h + \varepsilon) - u_h - \varepsilon) \end{aligned}$$

Again, since here $(u_{l,\varepsilon}, u_h + \varepsilon)$ is optimal for a seller with expertise \tilde{p}_l , we must have $\Pi(u_{l,\varepsilon}, u_h + \varepsilon; \tilde{p}_l) \geq \Pi(u_{l,\varepsilon}, u_h; \tilde{p}_l)$, which implies the following inequality,

$$\begin{aligned} & \tilde{p}_h(\Psi_h(u_h + \varepsilon) - \Psi_h(u_h))(S_h(u_{l,\varepsilon}, u_h + \varepsilon) - u_h - \varepsilon) \\ & \geq \tilde{p}_l \Psi_l(u_{l,\varepsilon})(S_l(u_{l,\varepsilon}, u_h) - S_l(u_{l,\varepsilon}, u_h + \varepsilon)) + \tilde{p}_h \Psi_h(u_h)(S_h(u_{l,\varepsilon}, u_h) - S_h(u_{l,\varepsilon}, u_h + \varepsilon) + \varepsilon) \end{aligned} \tag{C.9}$$

So, by (C.8) and (C.9), we can form a squeezed inequality for the derivative of $\Psi_h(\cdot)$,

$$\begin{aligned} & \frac{\frac{\tilde{p}_l}{\tilde{p}_h} \Psi_l(u_{l,\varepsilon})(S_l(u_{l,\varepsilon}, u_h) - S_l(u_{l,\varepsilon}, u_h + \varepsilon)) + \Psi_h(u_h)(S_h(u_{l,\varepsilon}, u_h) - S_h(u_{l,\varepsilon}, u_h + \varepsilon) + \varepsilon)}{\varepsilon(S_h(u_{l,\varepsilon}, u_h + \varepsilon) - u_h - \varepsilon)} \\ & \leq \frac{\Psi_h(u_h + \varepsilon) - \Psi_h(u_h)}{\varepsilon} \leq \\ & \frac{\frac{p_l}{p_h} \Psi_l(u_l)(S_l(u_l, u_h) - S_l(u_l, u_h + \varepsilon)) + \Psi_h(u_h)(S_h(u_l, u_h) - S_h(u_l, u_h + \varepsilon) + \varepsilon)}{\varepsilon(S_h(u_l, u_h + \varepsilon) - u_h - \varepsilon)} \end{aligned}$$

Considering the right-hand derivative, for $\varepsilon \searrow 0$ small $S_h(u_l, u_h + \varepsilon) - u_h - \varepsilon > 0$, since (u_l, u_h) is an equilibrium menu and (as we've argued in Section C in the proof ruling out atoms) $S_h(u_l, u_h) - u_h > 0$ for all such menus. Once the nonzero denominator is established, we can recall that $q_i(u_l, u_h)$ (specified in Theorem 4.1) is everywhere left as well as right differentiable in each variable and that $S_i(u_l, u_h) = \theta_i q_i(u_l, u_h) - \phi(q_i(u_l, u_h))$, so taking the limit of the right-hand expression, we obtain:

$$\frac{-\frac{p_l}{p_h} \Psi_l(u_l) \frac{\partial S_l}{\partial_+ u_h}(u_l, u_h) + \Psi_h(u_h) \left(1 - \frac{\partial S_h}{\partial_+ u_h}(u_l, u_h)\right)}{S_h(u_l, u_h) - u_h}$$

The left-hand side expression of the inequality is similar to the right, with the exception that $u_{l,\varepsilon}$ is substituted in for u_l and the weighting probabilities are $(\tilde{p}_l, \tilde{p}_h)$ instead of (p_l, p_h) , thus requiring a different argument. If $u_h - u_l > q_l^* \Delta\theta$, we establish in the orderedness proof that offerings locally around both u_l and u_h are those of dually efficient bids and that they preserve expected profits from both low and high matches - $\Pi_i(u_{l,\varepsilon}, u_h + \varepsilon) = \Pi_i(u_{l,\varepsilon}, u_h)$ for $i \in \{l, h\}$. So, if a seller of type \tilde{p}_l finds $(u_{l,\varepsilon}, u_h + \varepsilon)$ optimal, then so would a seller of type p_l , yielding a limit of the left-hand side expression as above (replacing \tilde{p}_l with p_l for ε small). On the other hand, if $u_h - u_l < q_l^* \Delta\theta$, IC_h binds locally for all menus with $\varepsilon > 0$ small, so the strict ordering of profits by type requires $\tilde{p}_l \searrow p_l$. Lastly, if $u_h - u_l = q_l^* \Delta\theta$, then we must be able to select a subsequence with $u_h - u_{l,\varepsilon} > u_h - u_l$ or $< u_h - u_l$. In the former (latter) case, IC_h is slack (binds) at the menus of the subsequence, so the argument from the $u_h - u_l > q_l^* \Delta\theta$ ($u_h - u_l < q_l^* \Delta\theta$) case applies, and this is sufficient to obtain a convergent left-hand inequality.

Therefore, the limit as $\varepsilon \rightarrow 0$ of the left-hand and right-hand inequalities is identical, thus establishing the right-hand derivative claim:

$$\frac{d\Psi_h}{d_+ u_h}(u_h) = \frac{-\frac{p_l}{p_h} \Psi_l(u_l) \frac{\partial S_l(u_l, u_h)}{\partial_+ u_h} + \Psi_h(u_h) \left(1 - \frac{\partial S_h(u_l, u_h)}{\partial_+ u_h}\right)}{S_h(u_l, u_h) - u_h}$$

The left-hand derivative argument is analogous except that we consider the menu $(u_l, u_h - \varepsilon)$.

As stated at the start, the case of $\Psi_l(\cdot)$ is analogous, ultimately yielding,

$$\frac{d\Psi_l}{d u_l}(u_l) = \frac{-\frac{p_h}{p_l} \Psi_h(u_h) \frac{\partial S_h(u_l, u_h)}{\partial u_l} + \Psi_l(u_l) \left(1 - \frac{\partial S_l(u_l, u_h)}{\partial u_l}\right)}{S_l(u_l, u_h) - u_l}$$

Under strictly convex costs, at the sole²⁴ points of possible non-differentiability where $u_h - u_l = \Delta\theta q_i^*$,

$$\frac{\partial S_i(u_l, u_h)}{\partial_+ u_i}(u_l, u_h) = \frac{\partial S_i(u_l, u_h)}{\partial_- u_i}(u_l, u_h) = \frac{\theta_i - \phi'(q_i^*)}{\Delta\theta} = 0$$

hence the stronger differentiability claim for Ψ_i . If costs instead take a piecewise form, unless $u_h - u_l = \Delta\theta q_i^*$, then $S_i(u_l, u_h)$ is locally linear (or constant) in each variable, so that the left and right

²⁴For a given equilibrium distribution.

derivatives are also equal. Since $S_i(u_l, u_h) - u_i$ are continuous in (u_l, u_h) and $\Psi_i(\cdot)$ are continuous by the lack of atoms, the distributions Ψ_i are also continuously differentiable on the interior of the supports with the exception of points u_i corresponding to the (no more than two) bids (u_l, u_h) at which $u_h - u_l = \Delta\theta q_i^*$. \square

References

- [1] Acquisti, A., L. Brandimarte, and G. Loewenstein (2015): “Privacy and Human Behavior in the Age of Information,” *Science*, 347(6221), 509-514.
- [2] Ali, S., G. Lewis, S. Vasserman (2023): “Voluntary Disclosure and Personalized Pricing,” *Review of Economics Studies*, 90(2), 538-571.
- [3] Babina, T. Fedyk, A. He, and J. Hodson (2022): “Artificial Intelligence, Firm Growth, and Product Innovation,” <https://ssrn.com/abstract=3651052>.
- [4] Bergeman, D., and A. Bonatti (2019): “Markets for Information: An Introduction,” *Annual Review of Economics*, 11, 85-107.
- [5] Bergemann, D., B. Brooks, S. Morris (2015) : “The Limits of Price Discrimination,” *American Economic Review*, 105, 921–957.
- [6] Braghieri, L. (2019): “Targeted Advertising and Price Discrimination in Intermediated Online Markets,” Working Paper, <https://ssrn.com/abstract=3072692>.
- [7] Bresnahan, T. (2010): “General Purpose Technologies,” *Handbook of the Economics of Innovation*, 2(1), 761–91.
- [8] Bresnahan, T., and M. Trajtenberg (1995): “General Purpose Technologies ‘Engines of Growth’?” *Journal of Econometrics*, 65(1), 83–108.
- [9] Brynjolfsson, E., D. Rock, and C. Syverson (2019): “Artificial Intelligence and the Modern Productivity Paradox: A Clash of Expectations and Statistics,” in *The Economics of Artificial Intelligence: An Agenda*, pp. 23–60. University of Chicago Press, Chicago, IL.
- [10] Burdett, K., and K. Judd (1983): “Equilibrium Price Dispersion,” *Econometrica*, 51 (4) (Jul., 1983), 955-969.
- [11] Bughin, J., M. Chui, N. Henke, J. Manyika, T. Saleh, G. Sethupathy, and B. Wiseman (2016) “The Age of Analytics: Competing in a Data-Driven World,” McKinsey Global Institute.
- [12] Calvano, E., G. Calzolari, V. Denicolò, S. Pastorello. *Product Recommendations and Market Concentration*. Working Paper, 2021.
- [13] Cam, A., M. Chui, and B. Hall (2019): “2019 Global AI Survey,” McKinsey Global Institute.
- [14] Chen, L., A. Mislove, and C. Wilson (2016): “An Empirical Analysis of Algorithmic Pricing on Amazon Marketplace,” in *Proceedings of the 25th International Conference on World Wide Web*, pp. 1339–1349.
- [15] Cho, J., T. DeStefano, H. Kim, J. Paik (2021): “What determines AI adoption?,” Working Paper.
- [16] Cockburn, I. M., R. Henderson, and S. Stern (2019): “The Impact of Artificial Intelligence on Innovation: An Exploratory Analysis,” in *Economics of Artificial Intelligence: An Agenda*, ed. by A. Agrawal, J. Gans, and A. Goldfarb, pp. 115-146. University of Chicago Press, Chicago, IL.
- [17] Council of Economic Advisers (2015): “The Economics of Big Data and Differential Pricing”.
- [18] Eeckhout, J., and L. Veldkamp (2022): “Data and Markups: A Macro-Finance Perspective,” Working Paper.

- [19] European Commission (2012): “Impact assessment accompanying the document Regulation of the European Parliament and of the Council on the protection of individuals with regard to the processing of personal data and on the free movement of such data (General Data Protection Regulation) and Directive of the European Parliament and of the Council on the protection of individuals with regard to the processing of personal data by competent authorities for the purposes of prevention, investigation, detection or prosecution of criminal offences or the execution of criminal penalties, and the free movement of such data”, SEC(2012) 72 final
- [20] Farboodi, M., and L. Veldkamp (2022): “A Model of the Data Economy,” Working Paper.
- [21] Federal Trade Commission (2012): “Protecting Consumer Privacy in an Era of Rapid Change: Recommendations for businesses and policymakers,” Federal Trade Commission.
- [22] Federal Trade Commission (2013): “Device Fingerprinting: Opportunities for FTC Involvement,” Federal Trade Commission.
- [23] Federal Trade Commission (2014): “Data Brokers: A Call for Transparency and Accountability,” Federal Trade Commission.
- [24] Feng, Y., R. Gradwohl, J. Hartline, and A. Johnsen (2019): “Bias-Variance Games,” *arXiv preprint arXiv:1909.03618*.
- [25] Fudenberg, D. and J. M. Villas-Boas (2007): “Behavior-Based Price Discrimination and Customer Recognition,” in *Economics and Information Systems*, ed. by T. J. Hendershott, vol. 1. Oxford: Elsevier Science.
- [26] Fudenberg, D. and J. M. Villas-Boas (2012): “Price Discrimination in the Digital Economy,” in *Oxford Handbook of the Digital Economy*, ed. by M. Peitz and J. Waldfogel. Oxford University Press.
- [27] Garrett, D., R. Gomes, and L. Maestri (2019): “Competitive Screening Under Heterogeneous Information,” *Review of Economic Studies*, 86(4), 1590-1630.
- [28] Goldfarb, A., B. Taska, and F. Teodoridis (2019): “Could Machine Learning Be a General-Purpose Technology? Evidence from Online Job Postings,” Working Paper, <https://ssrn.com/abstract=3468822>.
- [29] Hannak, A., G. Soeller, D. Lazer, A. Mislove, and C. Wilson (2014): “Measuring Price Discrimination and Steering on E-commerce Web Sites,” in *IMC '14 Proceedings of the 2014 Conference on Internet Measurement Conference*, pp. 305-318.
- [30] Harris, M. and A. Raviv (1993): “Differences of Opinion Make a Horse Race,” *Review of Financial Studies*, 6(3), 473-506.
- [31] Harrison, J. and D. Kreps (1978): “Speculative Investor Behavior in a Stock Market with Heterogeneous Expectations,” *Quarterly Journal of Economics*, 92(2), 323-336.
- [32] He, K., F. Sandomirskiy, and O. Tamuz (2021): “Private Private Information,” Working Paper.
- [33] Hidir, S. and N. Vellodi (2018): “Personalization, Discrimination and Information Revelation,” *Journal of the European Economic Association*, 19(2), 1342-1363.
- [34] Hobijn, B. and B. Jovanovic (2001): “The Information-Technology Revolution and the Stock Market:Evidence,” *American Economic Review*, 91(5), 1203-1220.

- [35] Hogan, K.(2018): “Consumer Experience in the Retail Renaissance: How Leading Brands Build a Bedrock with Data” <https://tinyurl.com/6n5vd2ax>.
- [36] Holmes, T.J. (1989): “ The Effects of Third-Degree Price Discrimination in Oligopoly, ” *American Economic Review*, 79, 244–250.
- [37] Ichihashi, S. (2020): “Online Privacy and Information Disclosure by Consumers,” *American Economic Review*, 110(2), 569–595.
- [38] Ichihashi, S. and A. Smolin (2023): “Buyer Optimal Algorithmic Consumption, ” Working Paper.
- [39] Kandel, E. and N. Pearson (1995): “Differential Interpretation of Public Signals and Trade in Speculative Markets,” *Journal of Political Economy*, 103(4), 831-872.
- [40] Kehoe, P., B. Larsen, and E. Pastorino (2020): “Dynamic Competition in the Era of Big Data,” Working paper.
- [41] Lakshmikantham, V. and L. Srinivasa (1969): *Differential and Integral Inequalities: Theory and Applications*, Academic Press, New York, 1969.
- [42] Lester, B., A. Shourideh, V. Venkateswaran, and A. Zetlin-Jones (2019): “Screening and Adverse Selection in Frictional Markets,” *Journal of Political Economy*, 127(1), 338 - 377.
- [43] Liang, A. (2021): “Games of Incomplete Information Played By Statisticians,” Working Paper, *arXiv preprint arXiv:1910.07018* .
- [44] Maschler, M., E. Solan, and S. Zamir (2013): *Game Theory*, Cambridge, UK: Cambridge University Press.
- [45] McNabb, A. (1986): “Comparison Theorems for Differential Equations,” *Journal of Mathematical Analysis and Applications*, 119(1-2), 417-428.
- [46] Mikians, J., L. Gyarmati, V. Erramilli, and N. Laoutaris (2012): “Detecting price and search discrimination on the internet,” in *Proceedings of the 11th ACM Workshop on Hot Topics in Networks*, pp. 79–84.
- [47] Mikians, J., L.Gyarmati, V. Erramilli, and N. Laoutaris (2013): “Crowd-Assisted Search for Price Discrimination in E-commerce: First Results,” in *Proceedings of the 9th ACM Conference on Emerging Networking Experiments and Technologies*, pp. 1–6, New York, Association for Computing Machinery.
- [48] Mussa, M. and S. Rosen (1978): “Monopoly and Product Quality,” *Journal of Economic Theory*, 18(2), 301-317.
- [49] Olea, J., P. Ortoleva, M. Pai, and A. Prat (2022): “Competing Models,” *The Quarterly Journal of Economics*, 137(4), 2419-2457.
- [50] Ottaviani, M. and P. Sørensen (2015): “Price Reaction to Information with Heterogeneous Beliefs and Wealth Effects: Underreaction, momentum, and reversal,” *American Economic Review*, 105(1), 1–34.
- [51] Robinson, J. (1933). “The Economics of Imperfect Competition, ” Macmillan, London.
- [52] Rothschild, M. and J. Stiglitz (1976): “Equilibrium in Competitive Insurance Markets: An Essay on the Economics of Imperfect Information,” *The Quarterly Journal of Economics* 90(4), 629-649.

- [53] Schmalensee, R. (1981). “Output and Welfare Implications of Monopolistic Third-Degree Price Discrimination,” *American Economic Review*, 71(1): 242-247.
- [54] Scheinkman, J. and W. Xiong (2013): “Overconfidence and Speculative Bubbles,” *Journal of Political Economy*, 111(6), 1183–1220.
- [55] Stole, L. (2007): “Price Discrimination and Competition,” In M. Armstrong and R. Porter eds., *The Handbook of Industrial Organization*, Elsevier, Amsterdam.
- [56] Trajtenberg, M. (2019): “Artificial Intelligence as the Next GPT,” in *Economics of Artificial Intelligence: An Agenda*, ed. by A. Agrawal, J. Gans, and A. Goldfarb, pp. 175-186. University of Chicago Press, Chicago, IL.
- [57] Turow, J., L. Feldman, and K. Meltzer (2005): “Open to Exploitation: America’s Shoppers Online and Offline,” Annenberg Public Policy Center.
- [58] UK Competition & Markets Authority (2021): “Algorithms: How they can reduce competition and harm consumers,” <https://www.gov.uk/government/publications/algorithms-how-they-can-reduce-competition-and-harm-consumers>
- [59] Varian, H. R. (1985): “Price Discrimination and Social Welfare ,” *American Economic Review*, 75(4): 870-875.
- [60] Yang, K.H. (2022): “Selling Consumer Data for Profit: Optimal Market Segmentation Design and its Consequences,” *American Economic Review*, 112(4), 1364-1393.